

Teasing Out Causal Relationships in Distributed Systems by Adding Delays to gRPCs

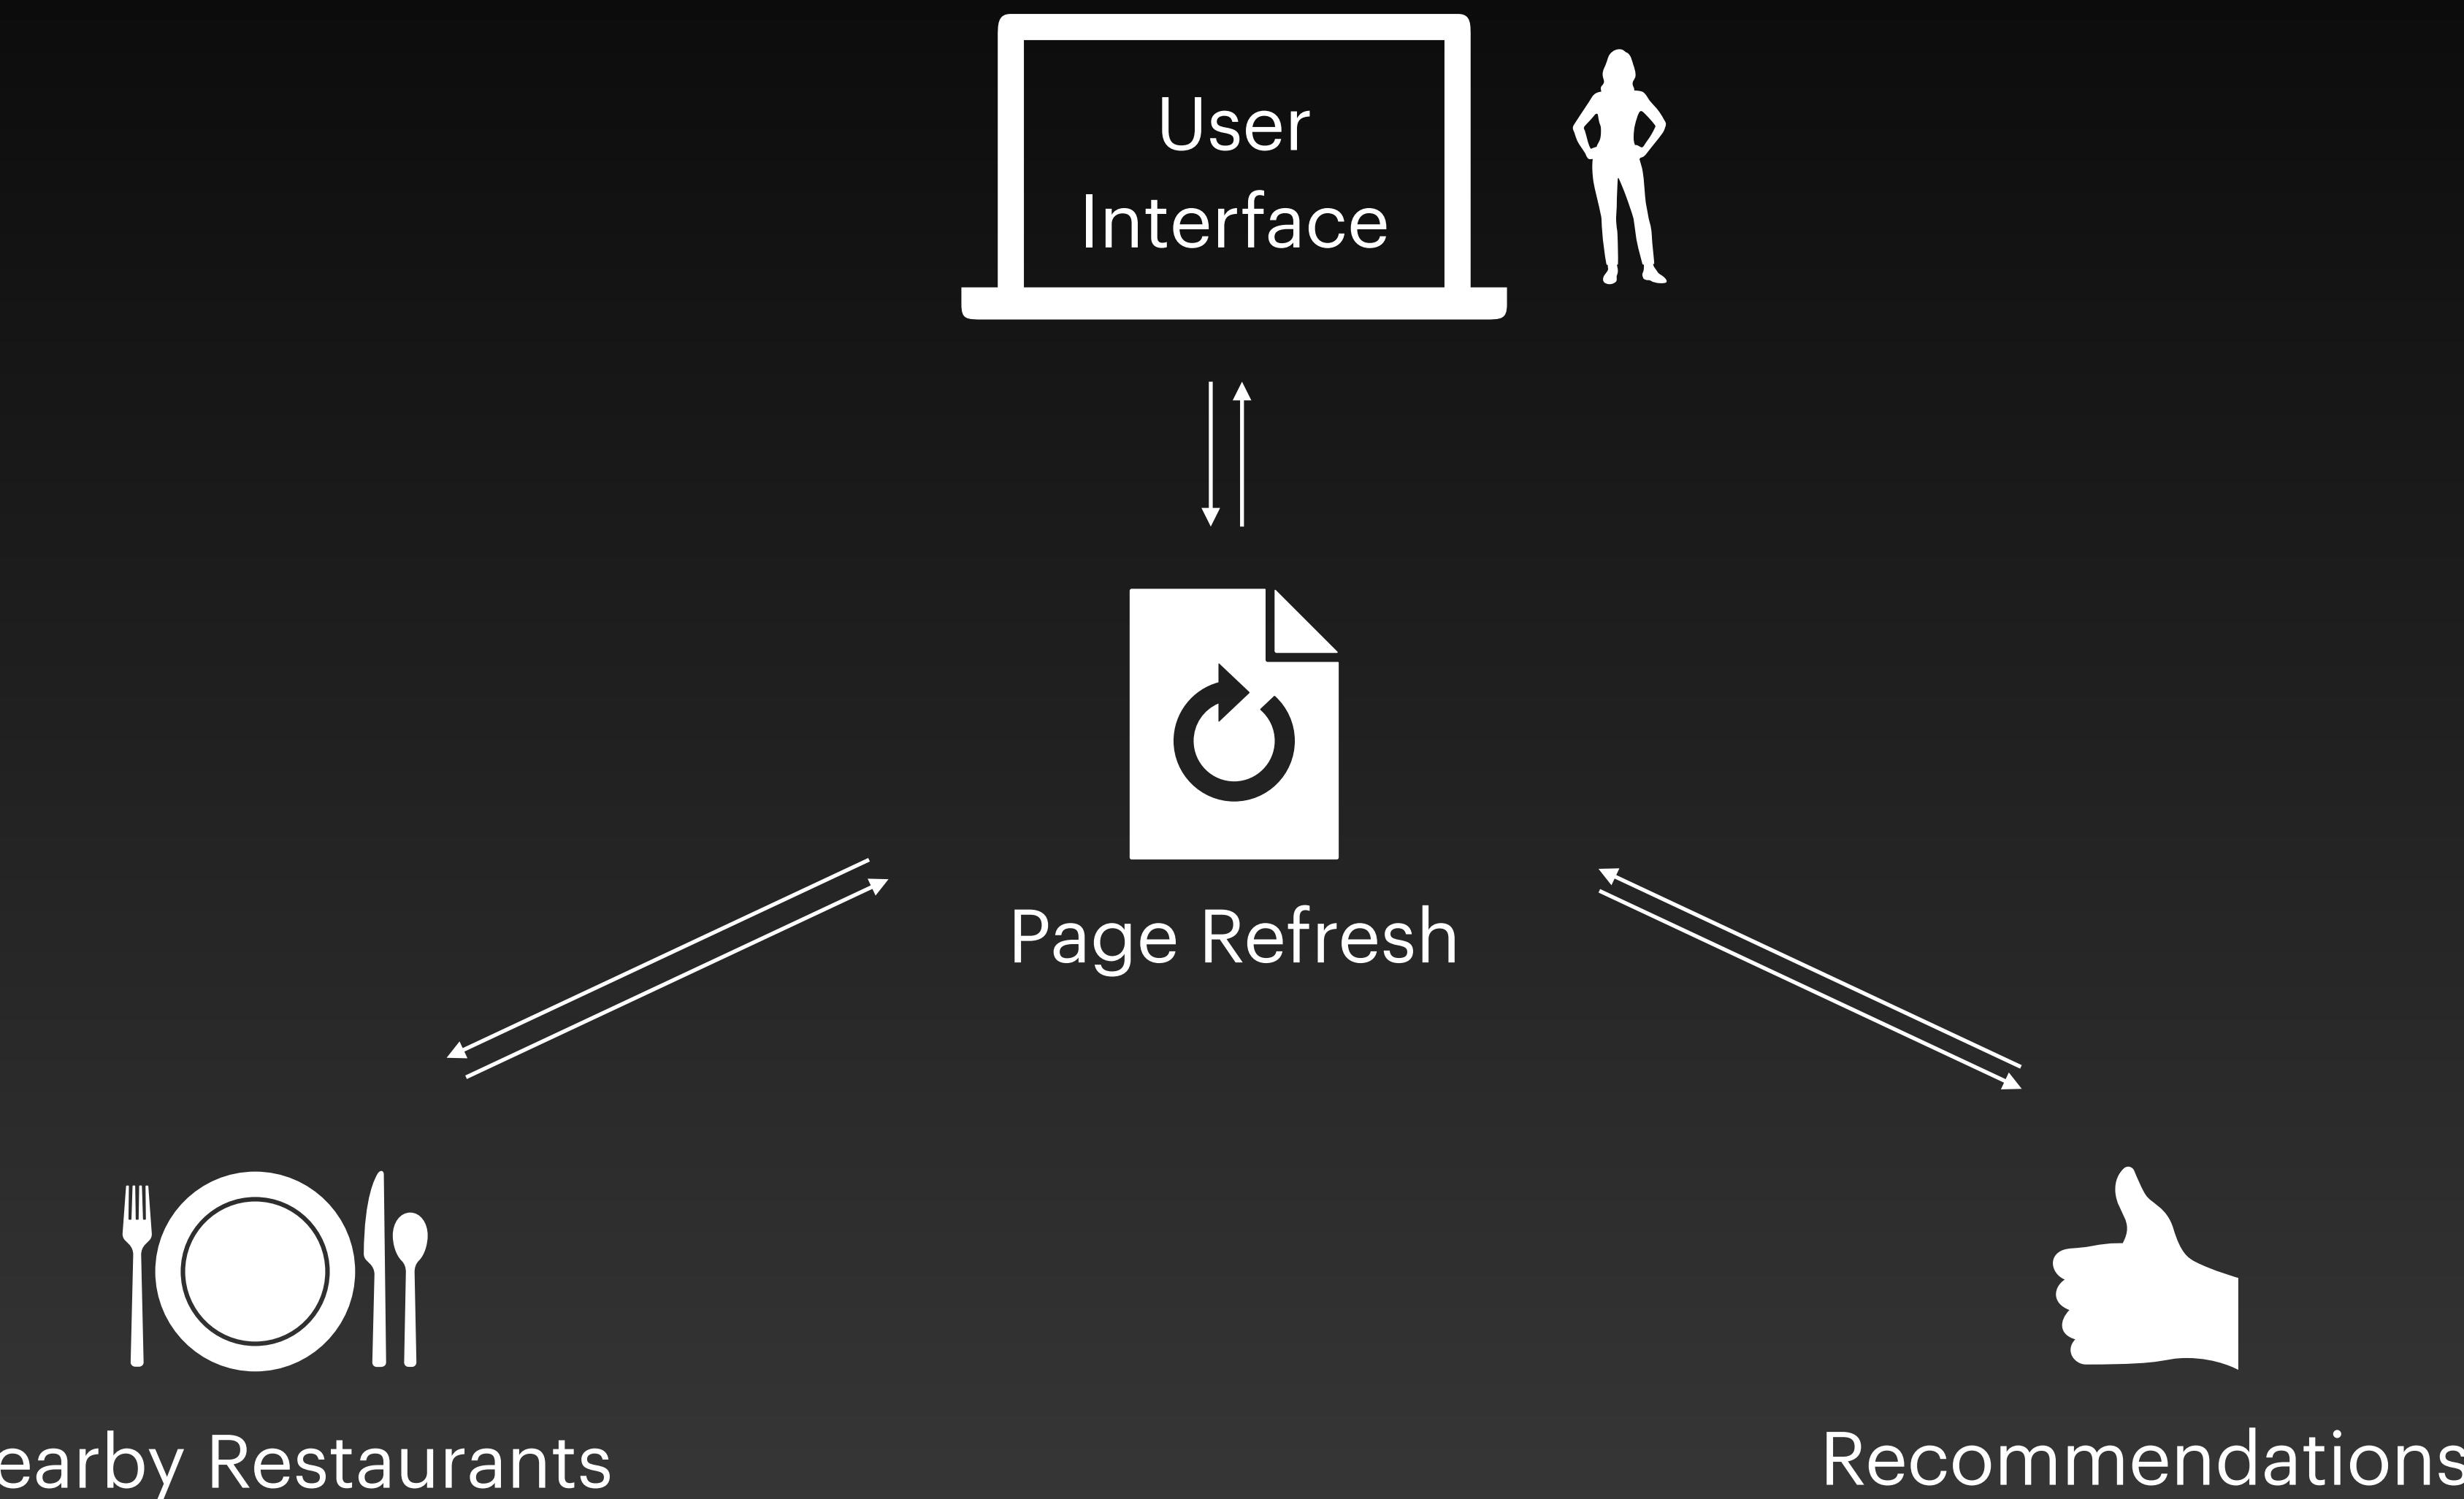
Mona Ma

May 10th 2024

Introduction

- Definitions
- Motivation

Microservices - Parent & Siblings



Execution Design

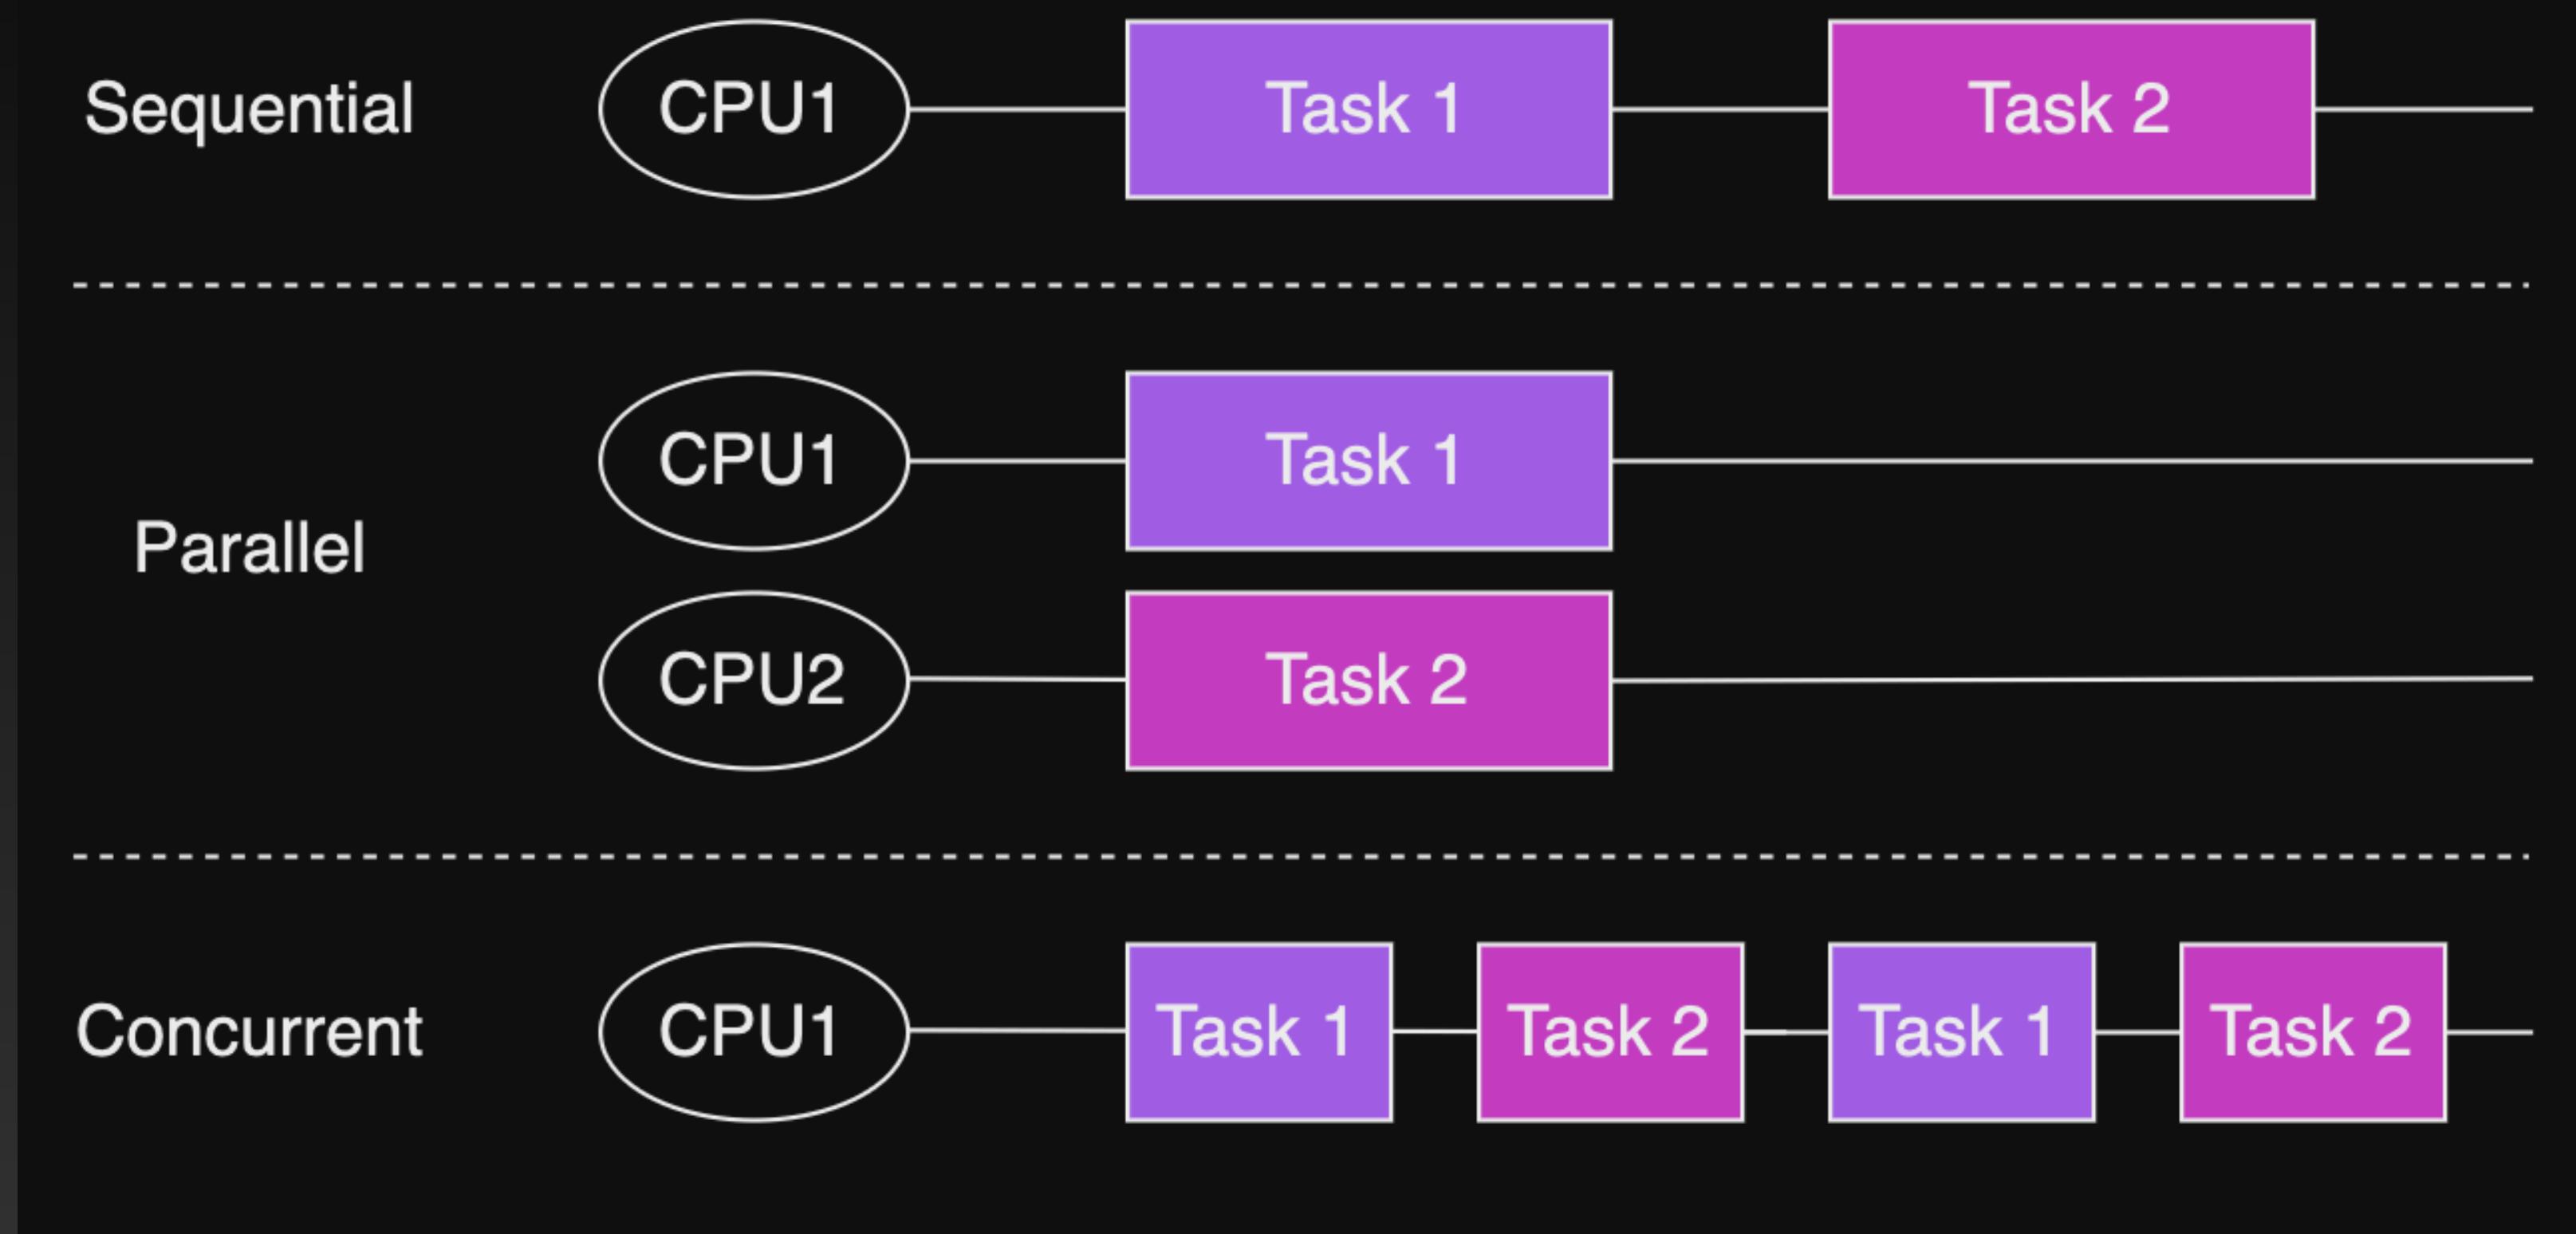


Figure 1: Visualization of Workflow

Challenge

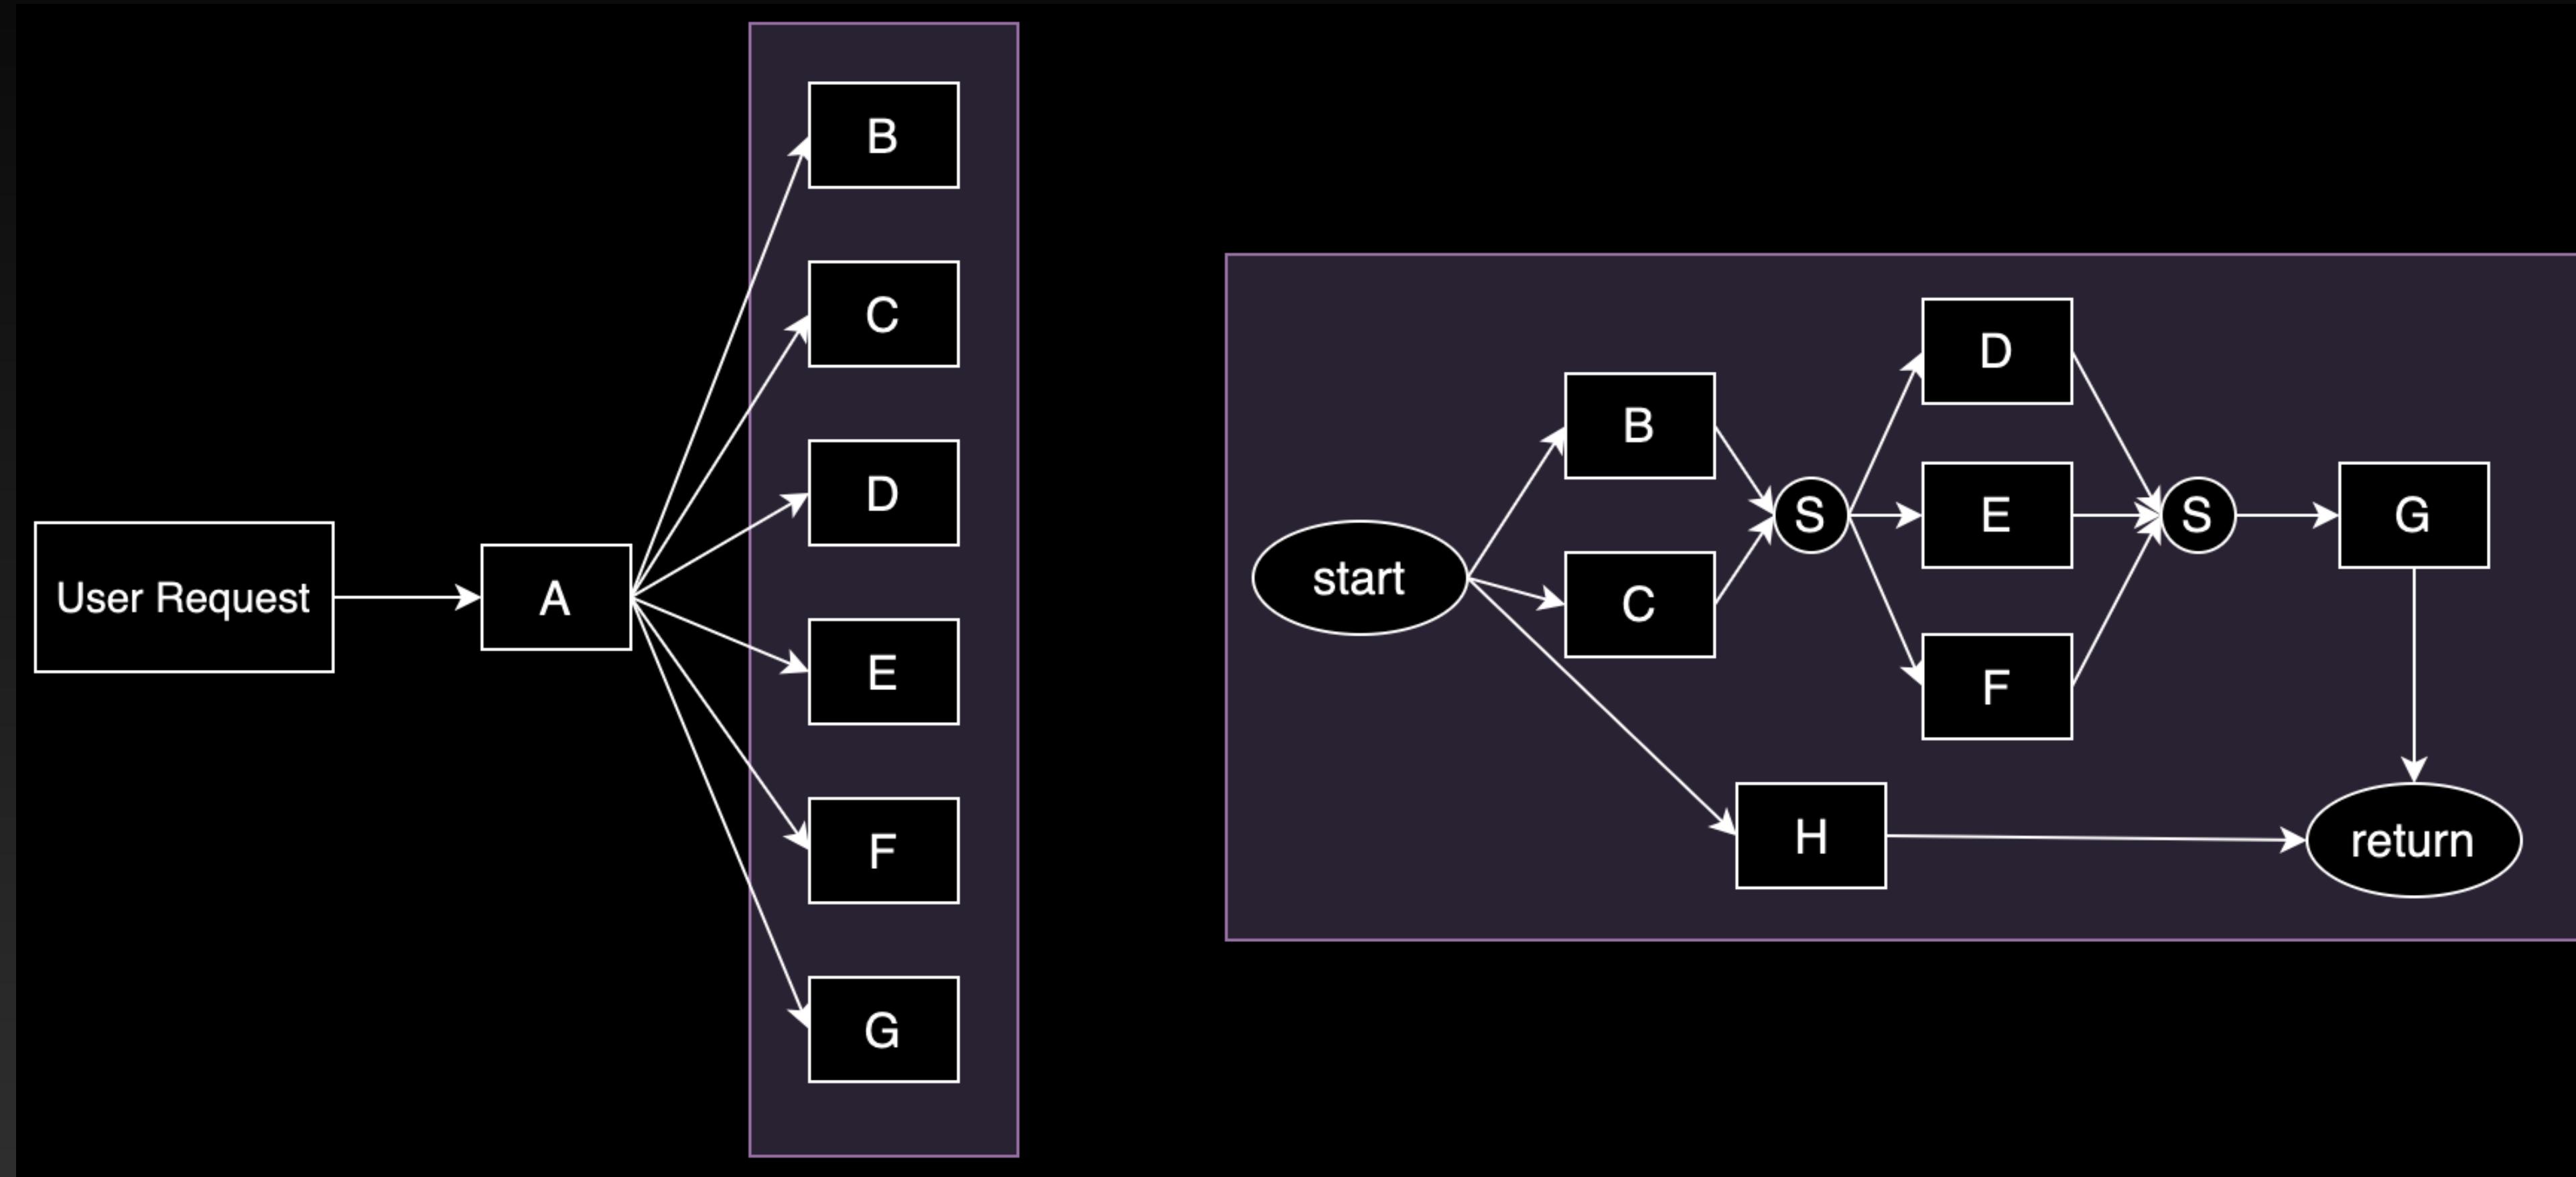


Figure 2: A Black Box within the System

Caller/Callee

Versus

Dependency

Problem statement

- Conduct a dependency analysis
- **Goal:** determine whether the siblings of a parent service are executed sequentially or in parallel
- **Outcome:** what can be parallelized -> performance optimization

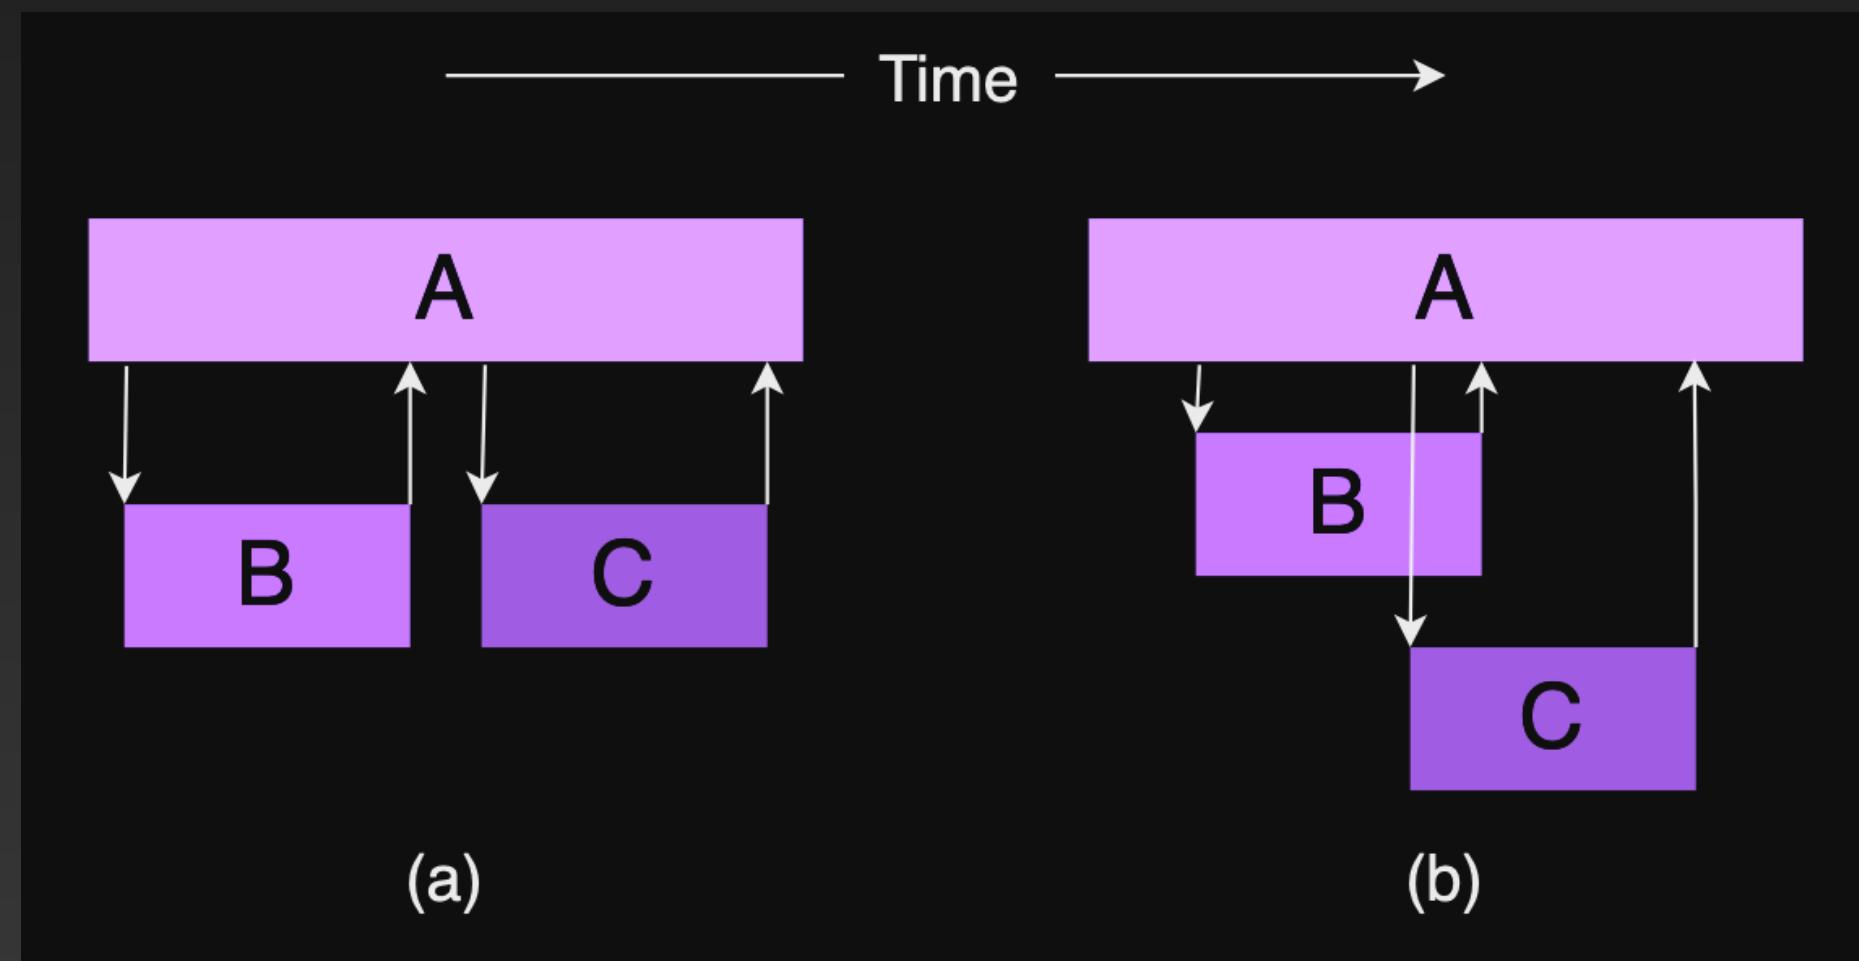


Figure 3: Sibling Services in a Simple Scenario

Deferred/Asynchronous Computation

Programming Languages	Functionality
Python	concurrent.futures async-io
JavaScript	promises, async/await
Java	java.util.concurrent, CompletableFuture
Go	Goroutines, Channels
...	...

Table 1: Different Programming Languages and Their Related Functions for Implementing Concurrency

Approach

- Active perturbation
- Add delays to selective sibling's return gRPC call

Mechanism

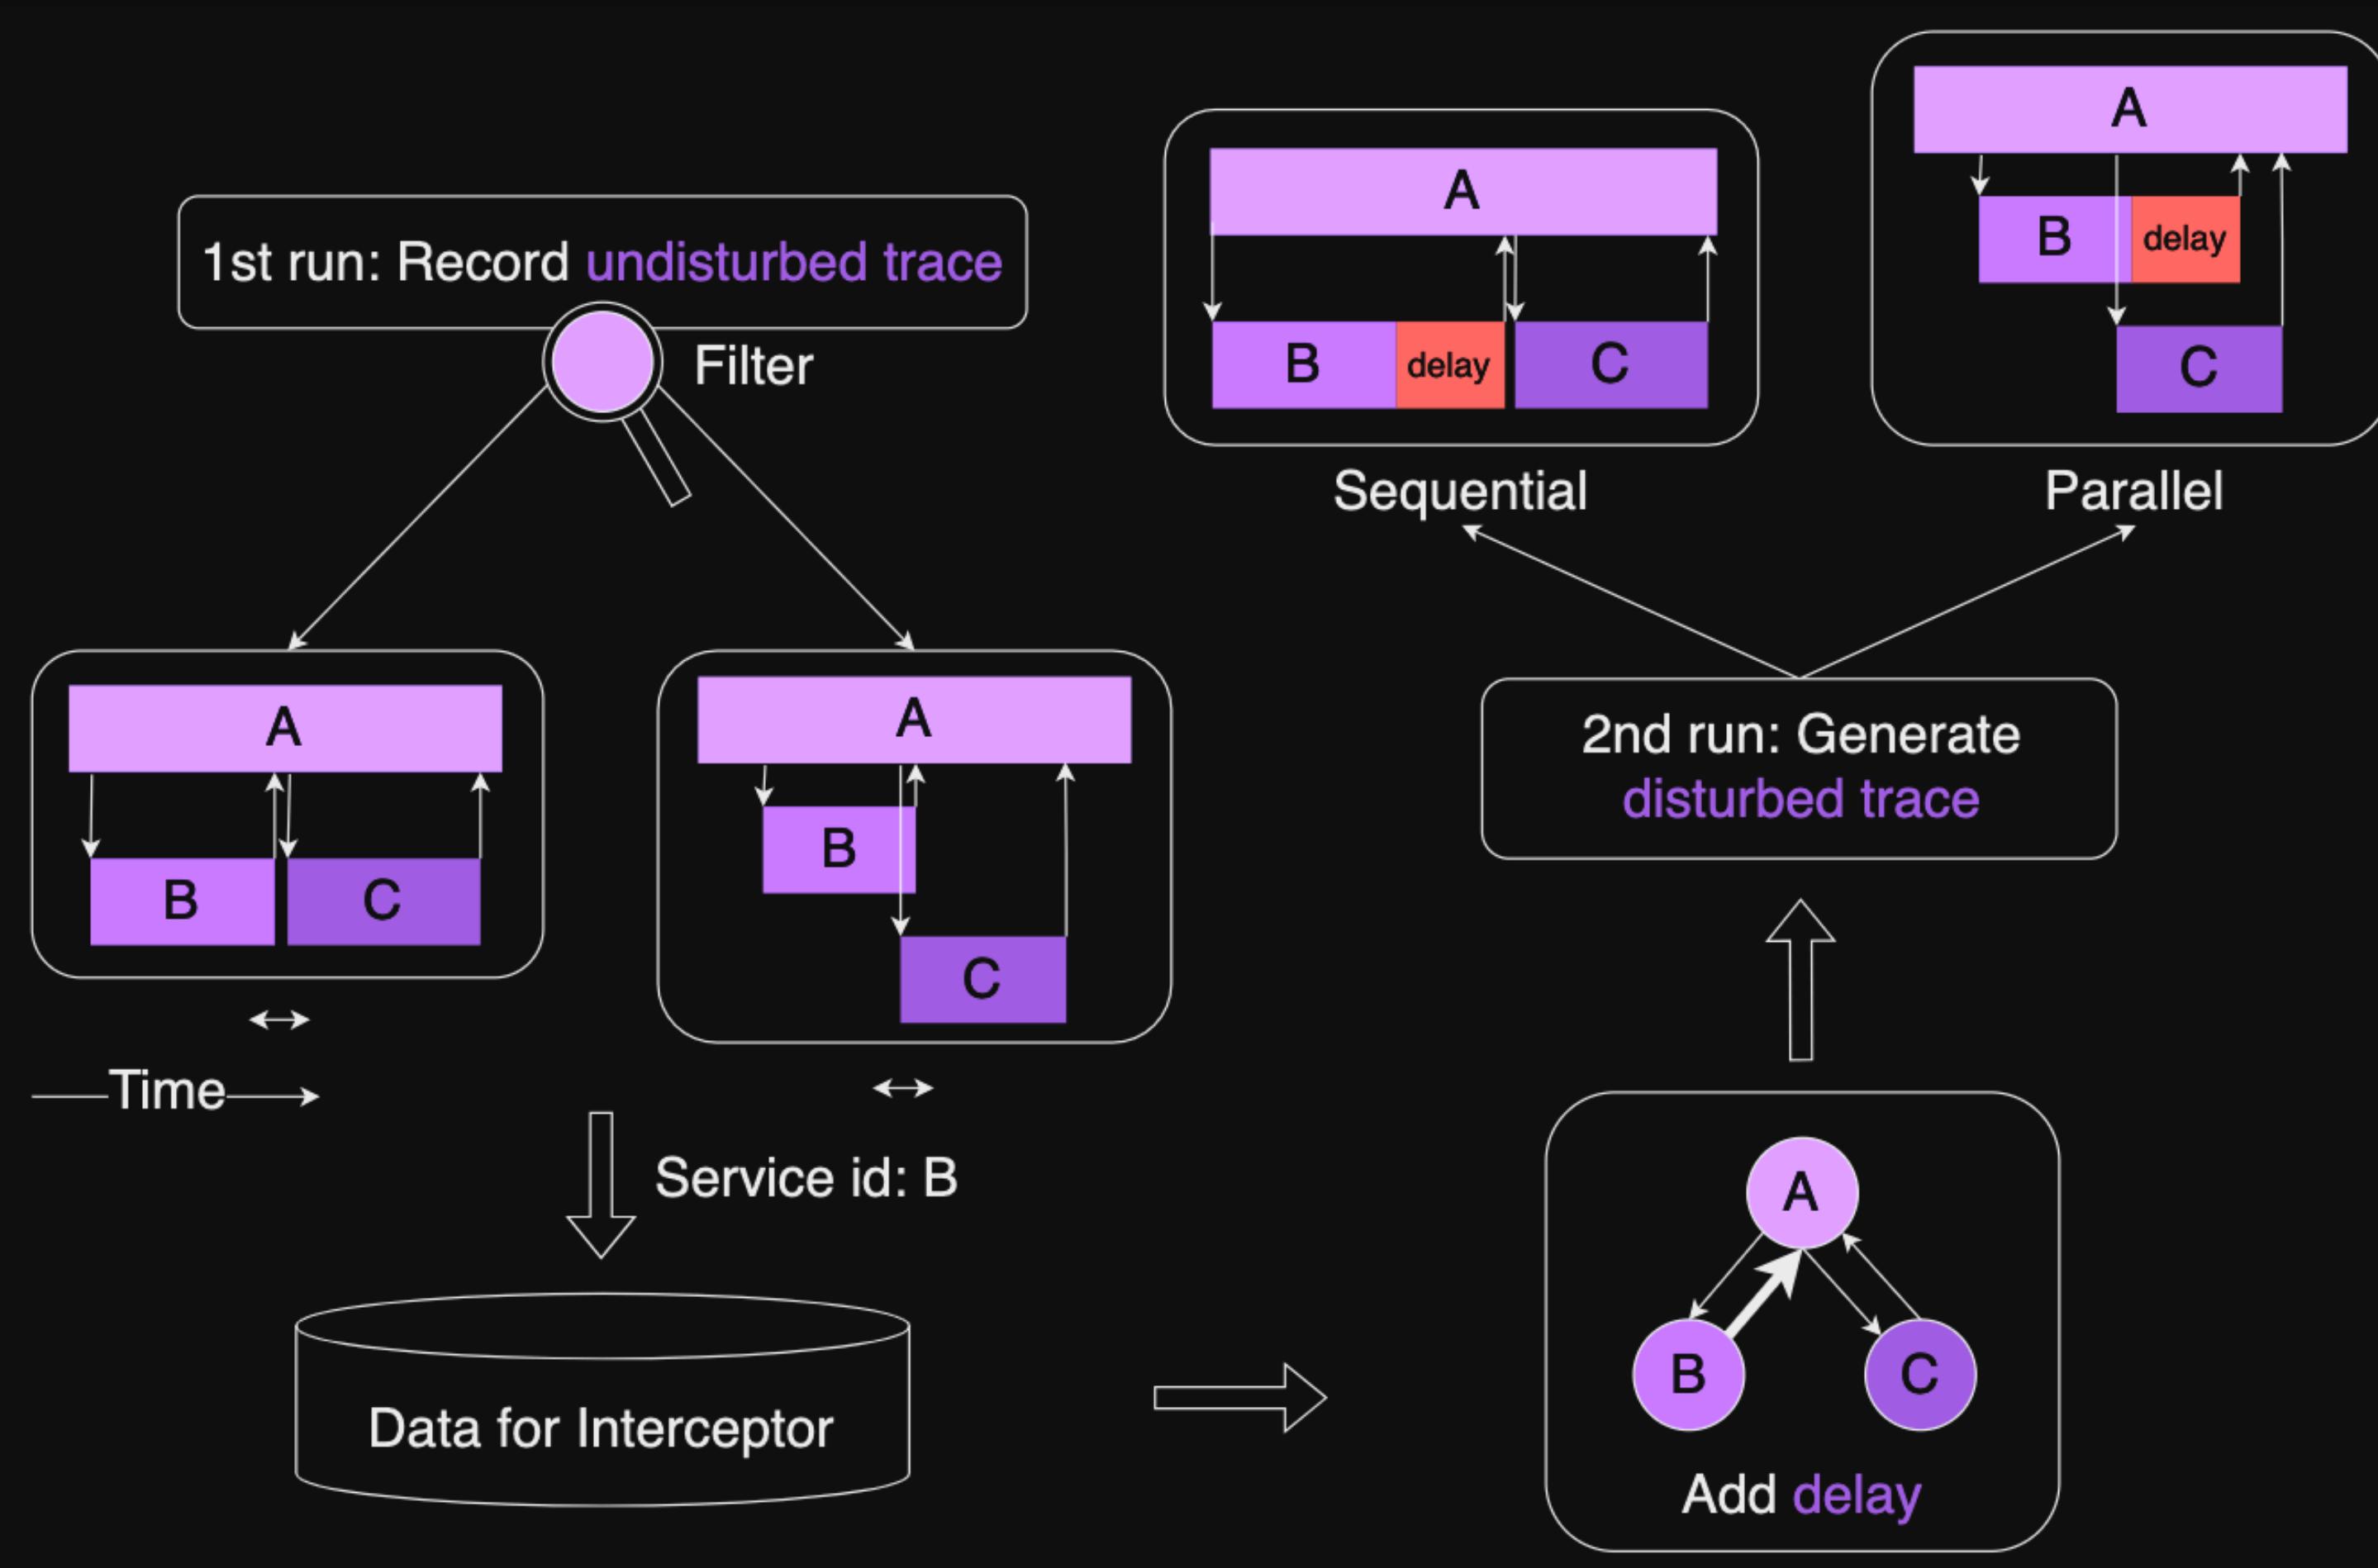


Figure 4: Project Workflow

- 2 Executions in total
- 1st Run:
 - Identify locality
 - Record interested services & timestamps
- 2nd Run:
 - Delay specific retuned response
 - Record timestamps
- Compare the results to summarize happen-before relationships

RPC & gRPC

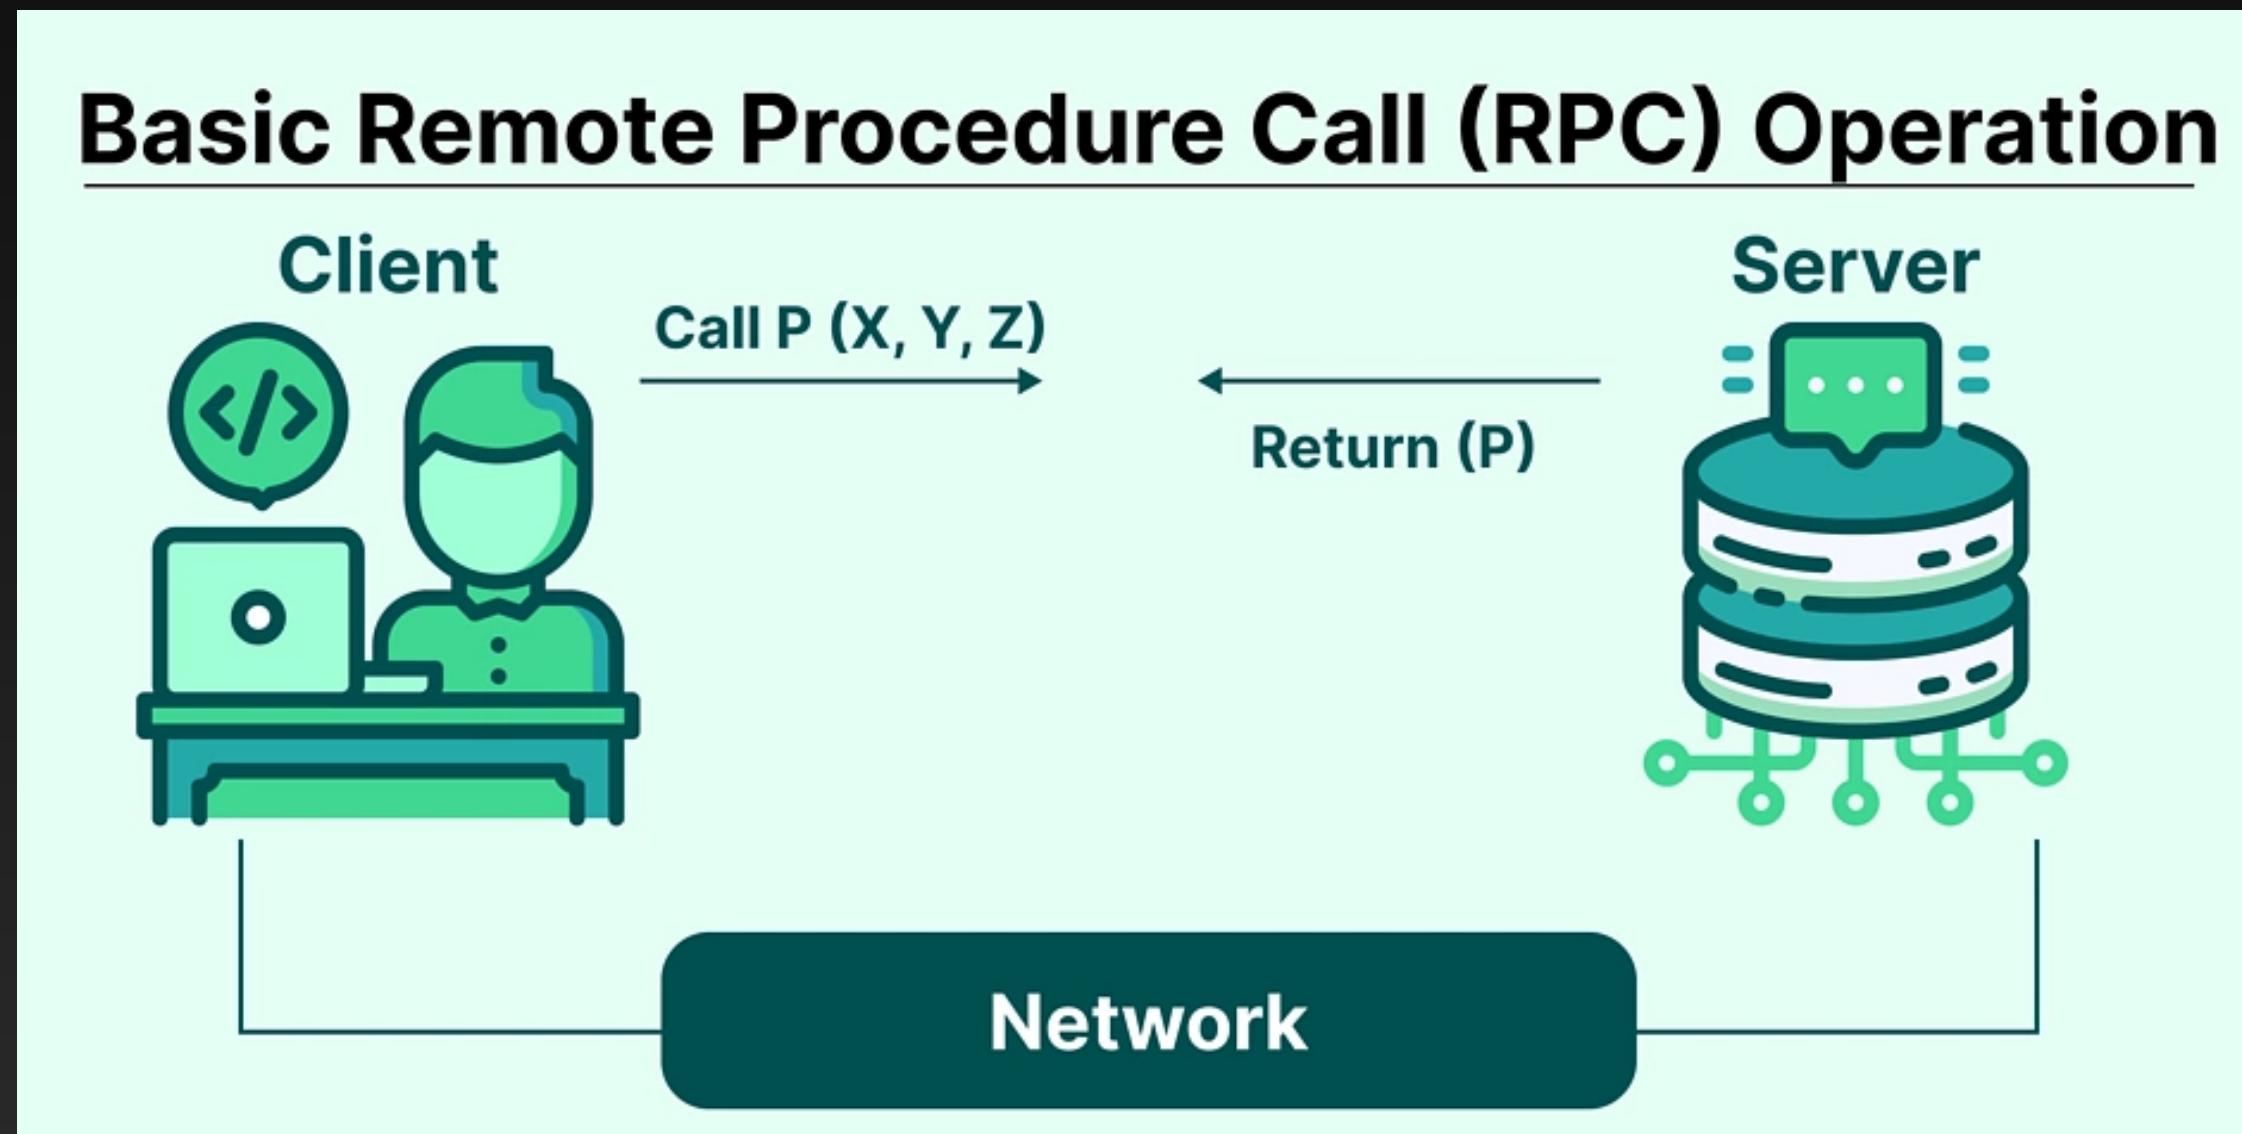


Figure 5: RPC

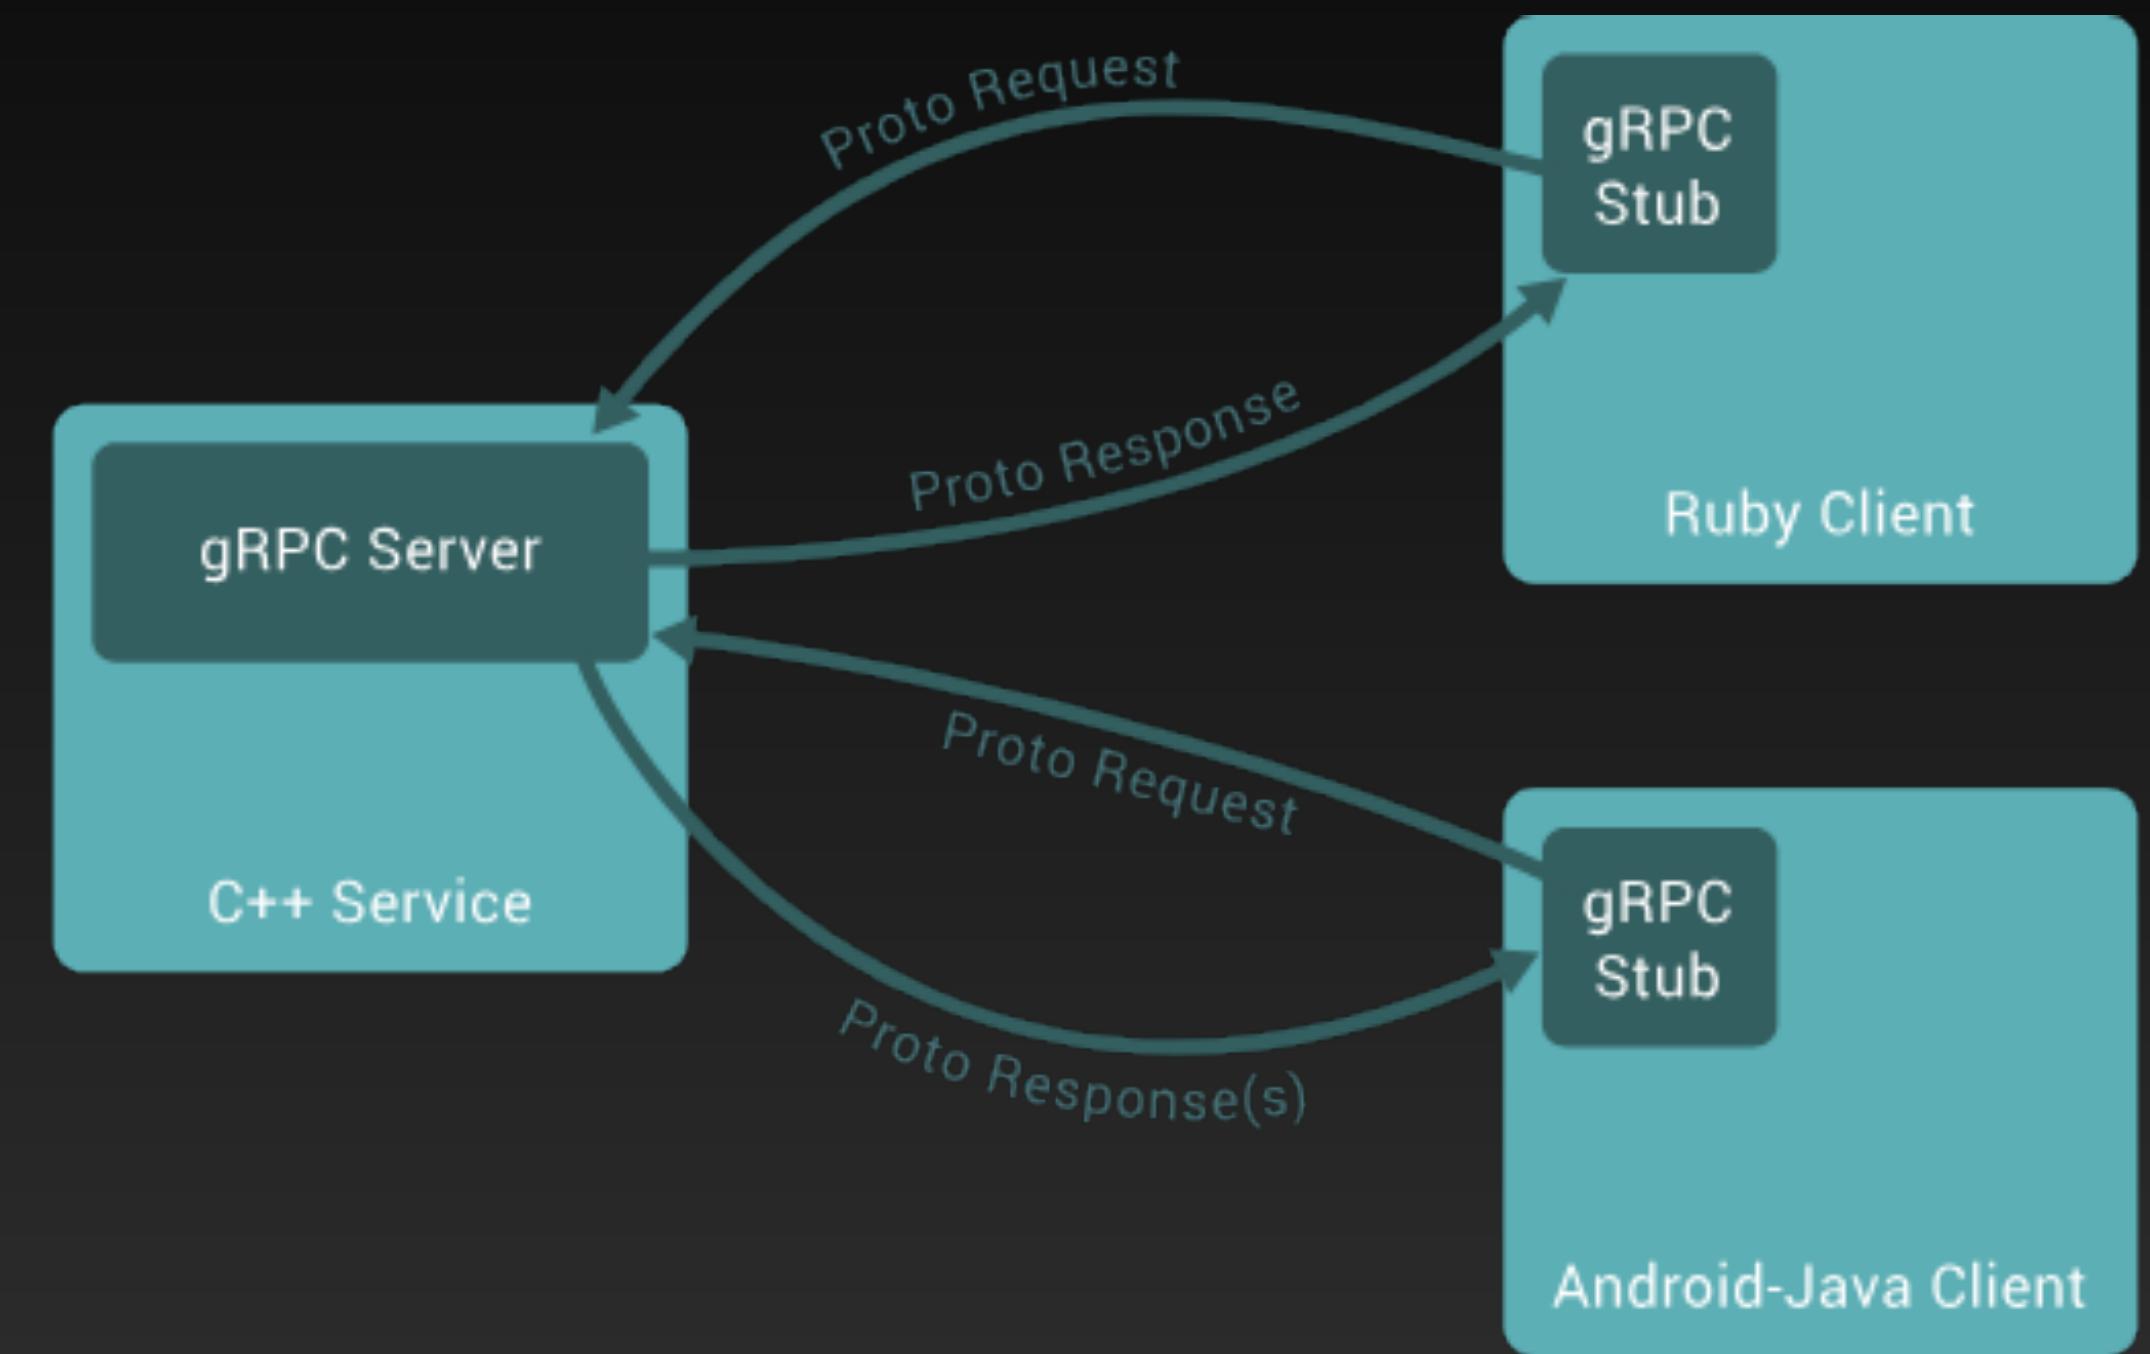


Figure 6: gRPC

gRPC Interceptor

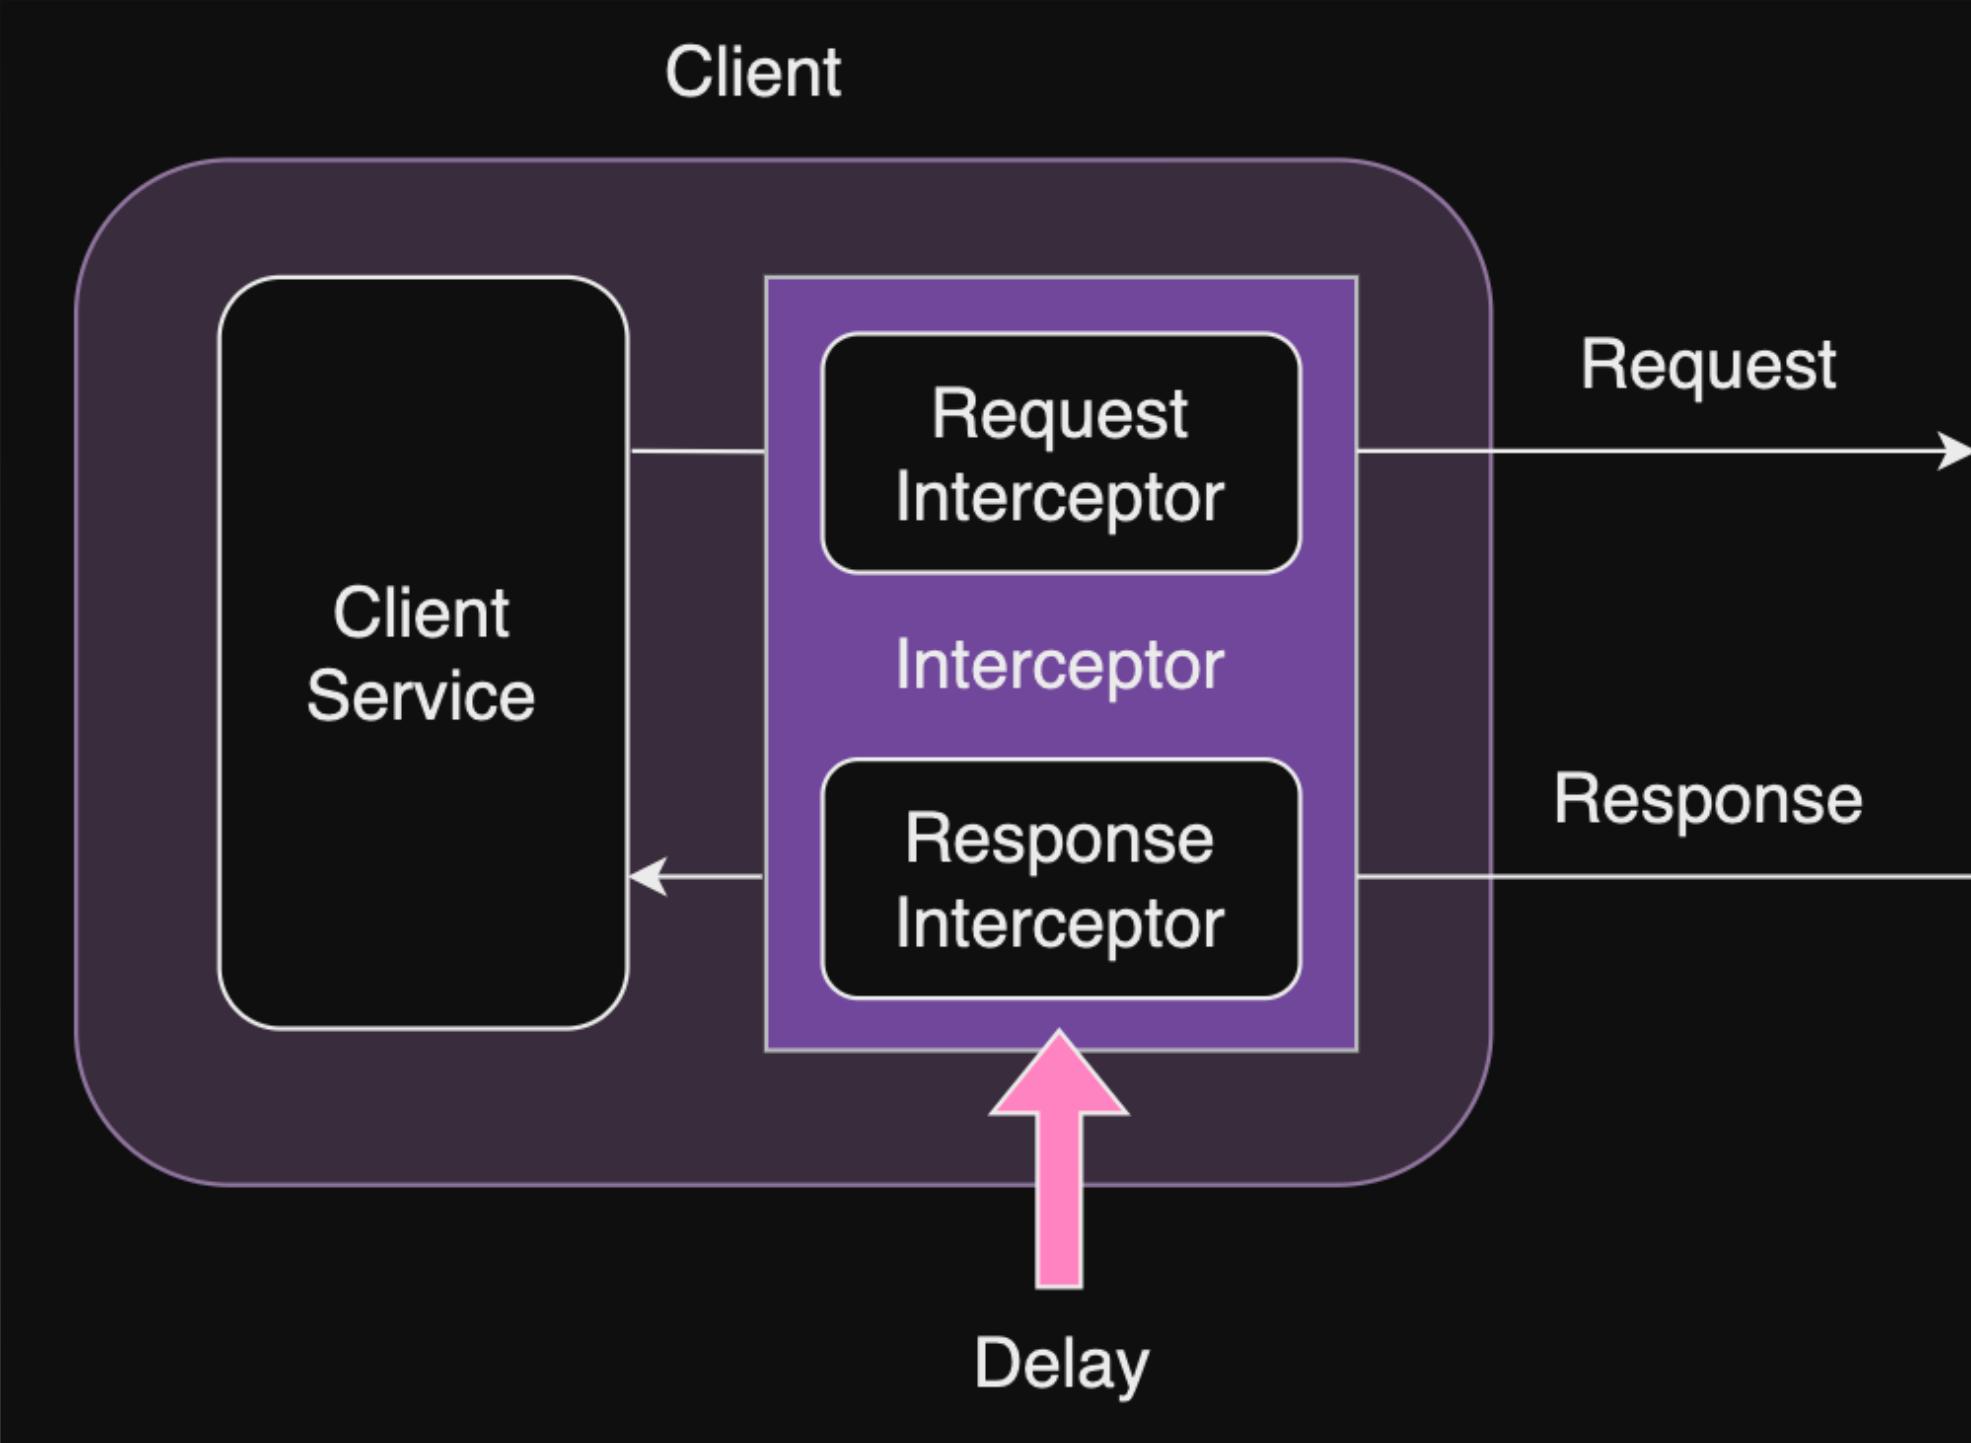


Figure 7: Client Side Interceptors

- Filters/middleware
- Popular use cases:
 - Metadata handling
 - Logging
 - Authentication & Authorization
 - Add delays (e.g. 0.001 seconds)

Implementation

- Reusable gRPC interceptor class & functions
- Input parameters:
 - `delay_ms: int`
 - `service_ids: Optional[list[str]]`
- Output:
 - Timestamps
- Toy Demo: Total 270 lines of code
- Interceptor: 15 lines of code

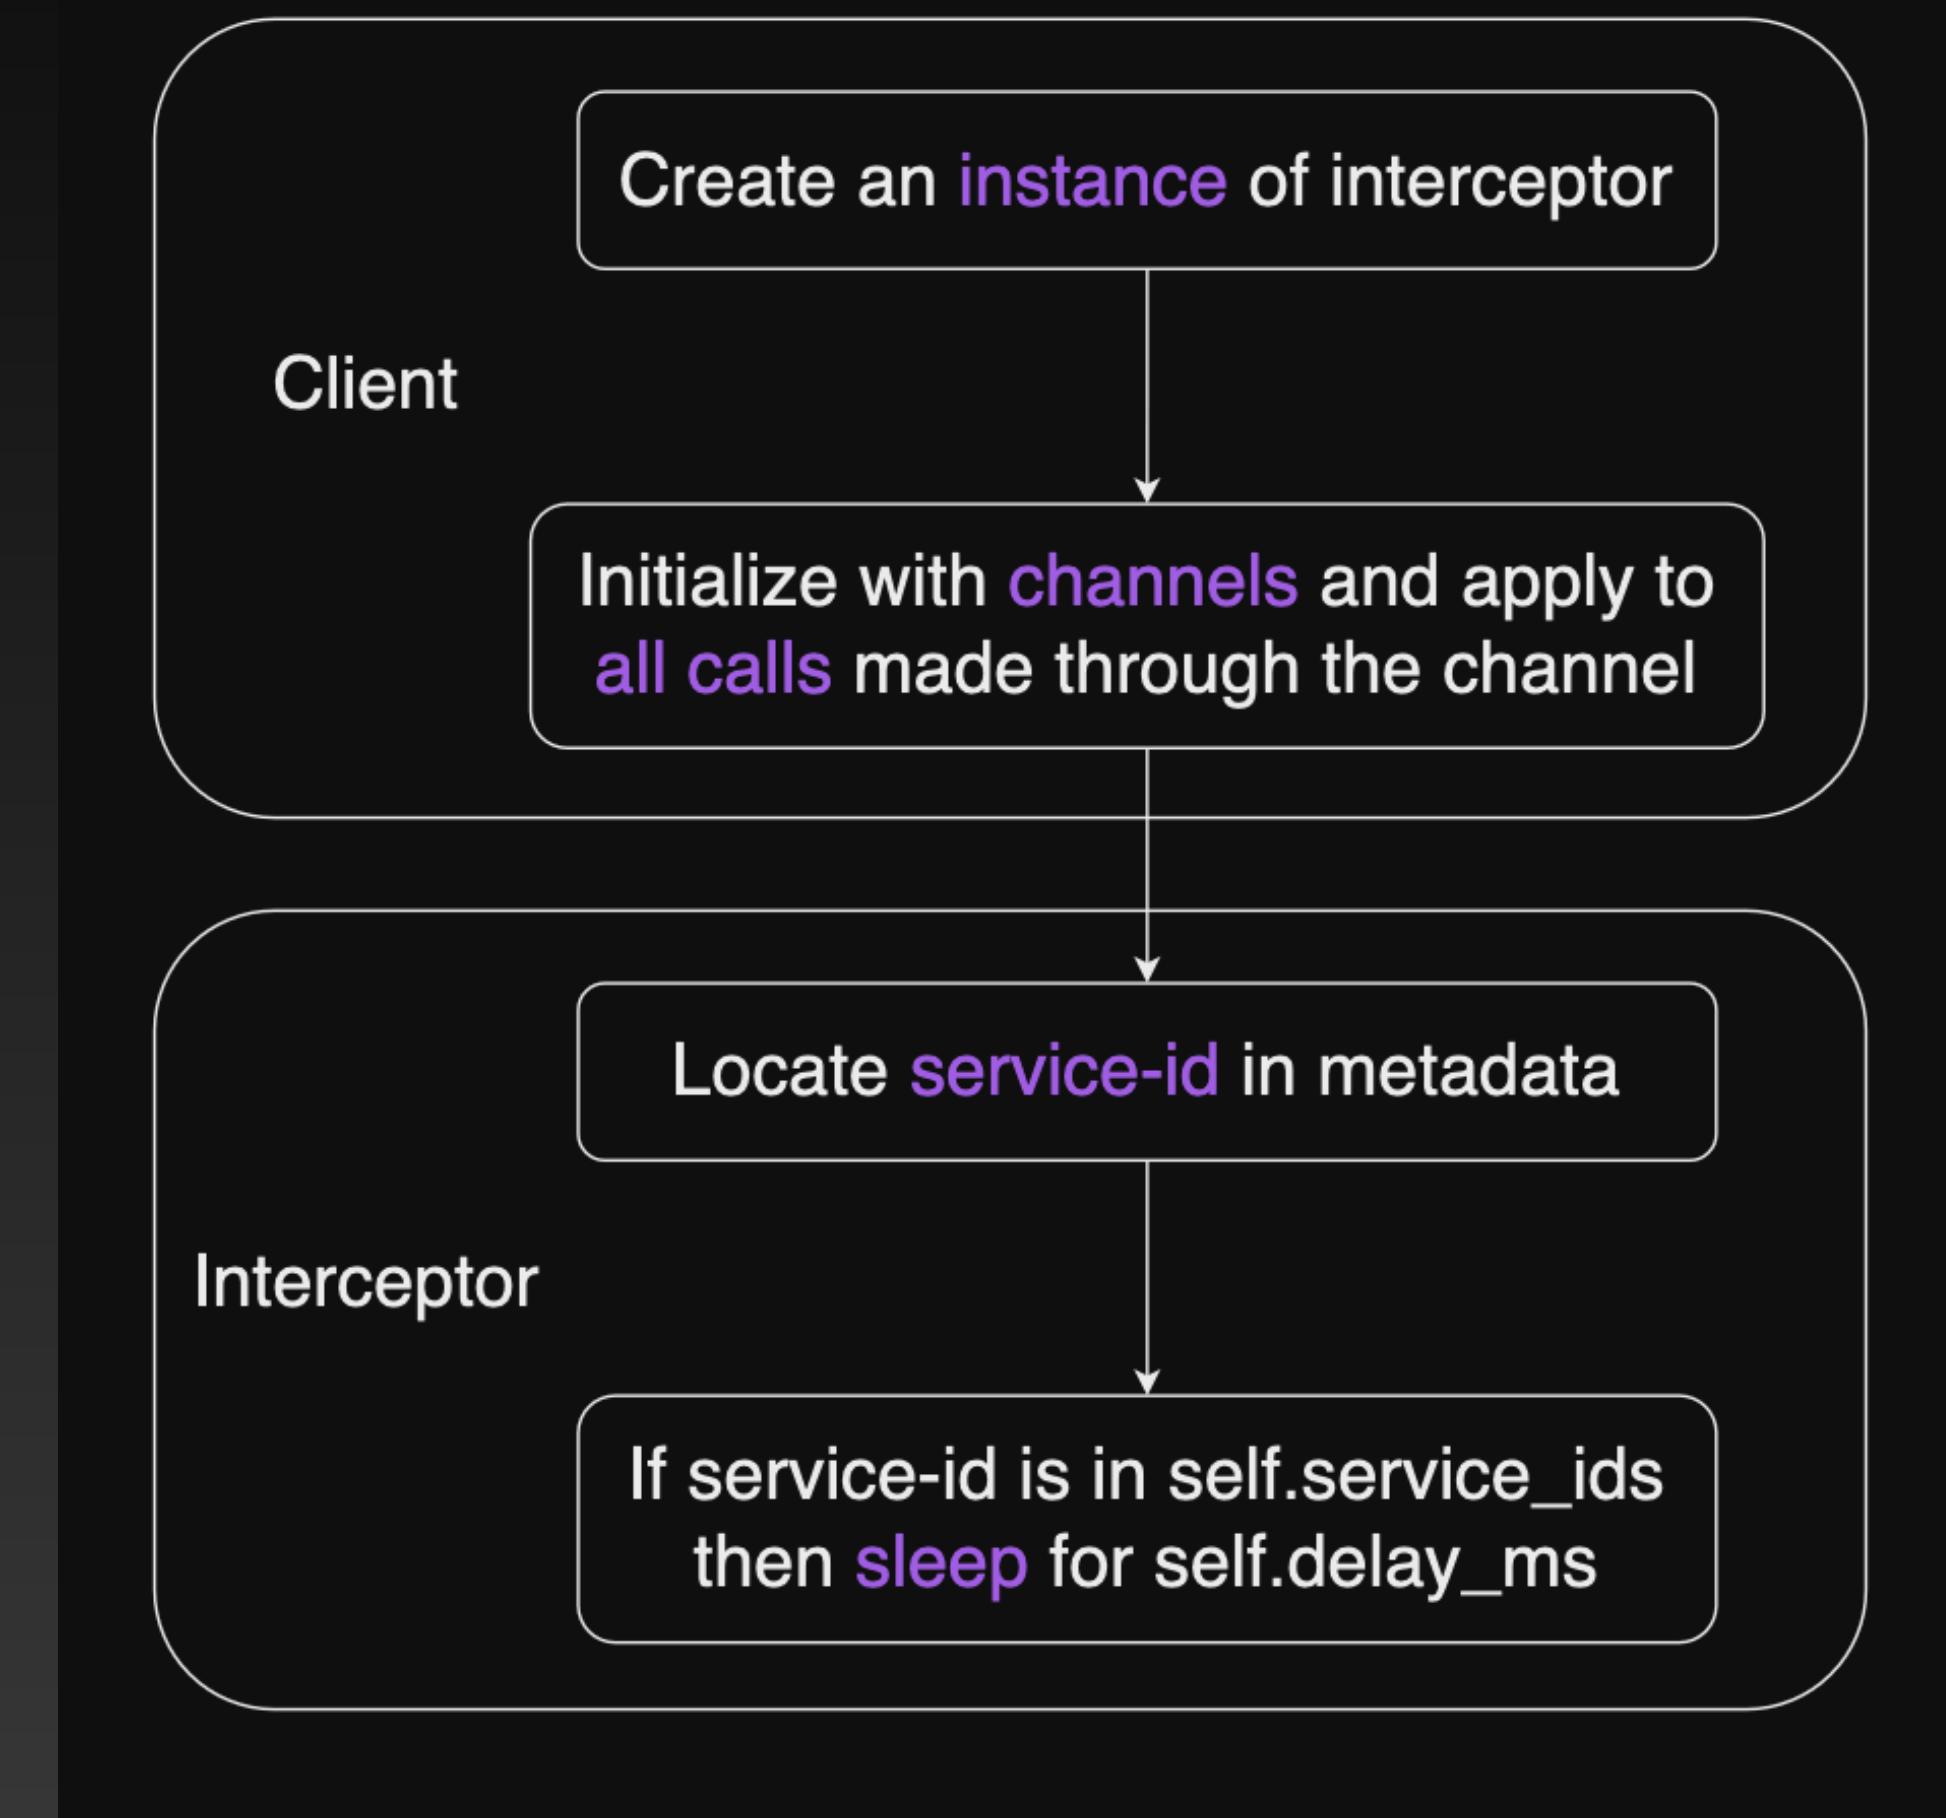


Figure 8: Implementation Steps

Toy Demo

- Using Python `async.io`
- 2 Microservices
 - Restaurant
 - Recommendation
- 4 gRPC calls
 - 2 requests & 2 responses
- 3 Modes
 - Concurrent
 - Sequential
 - Multiple

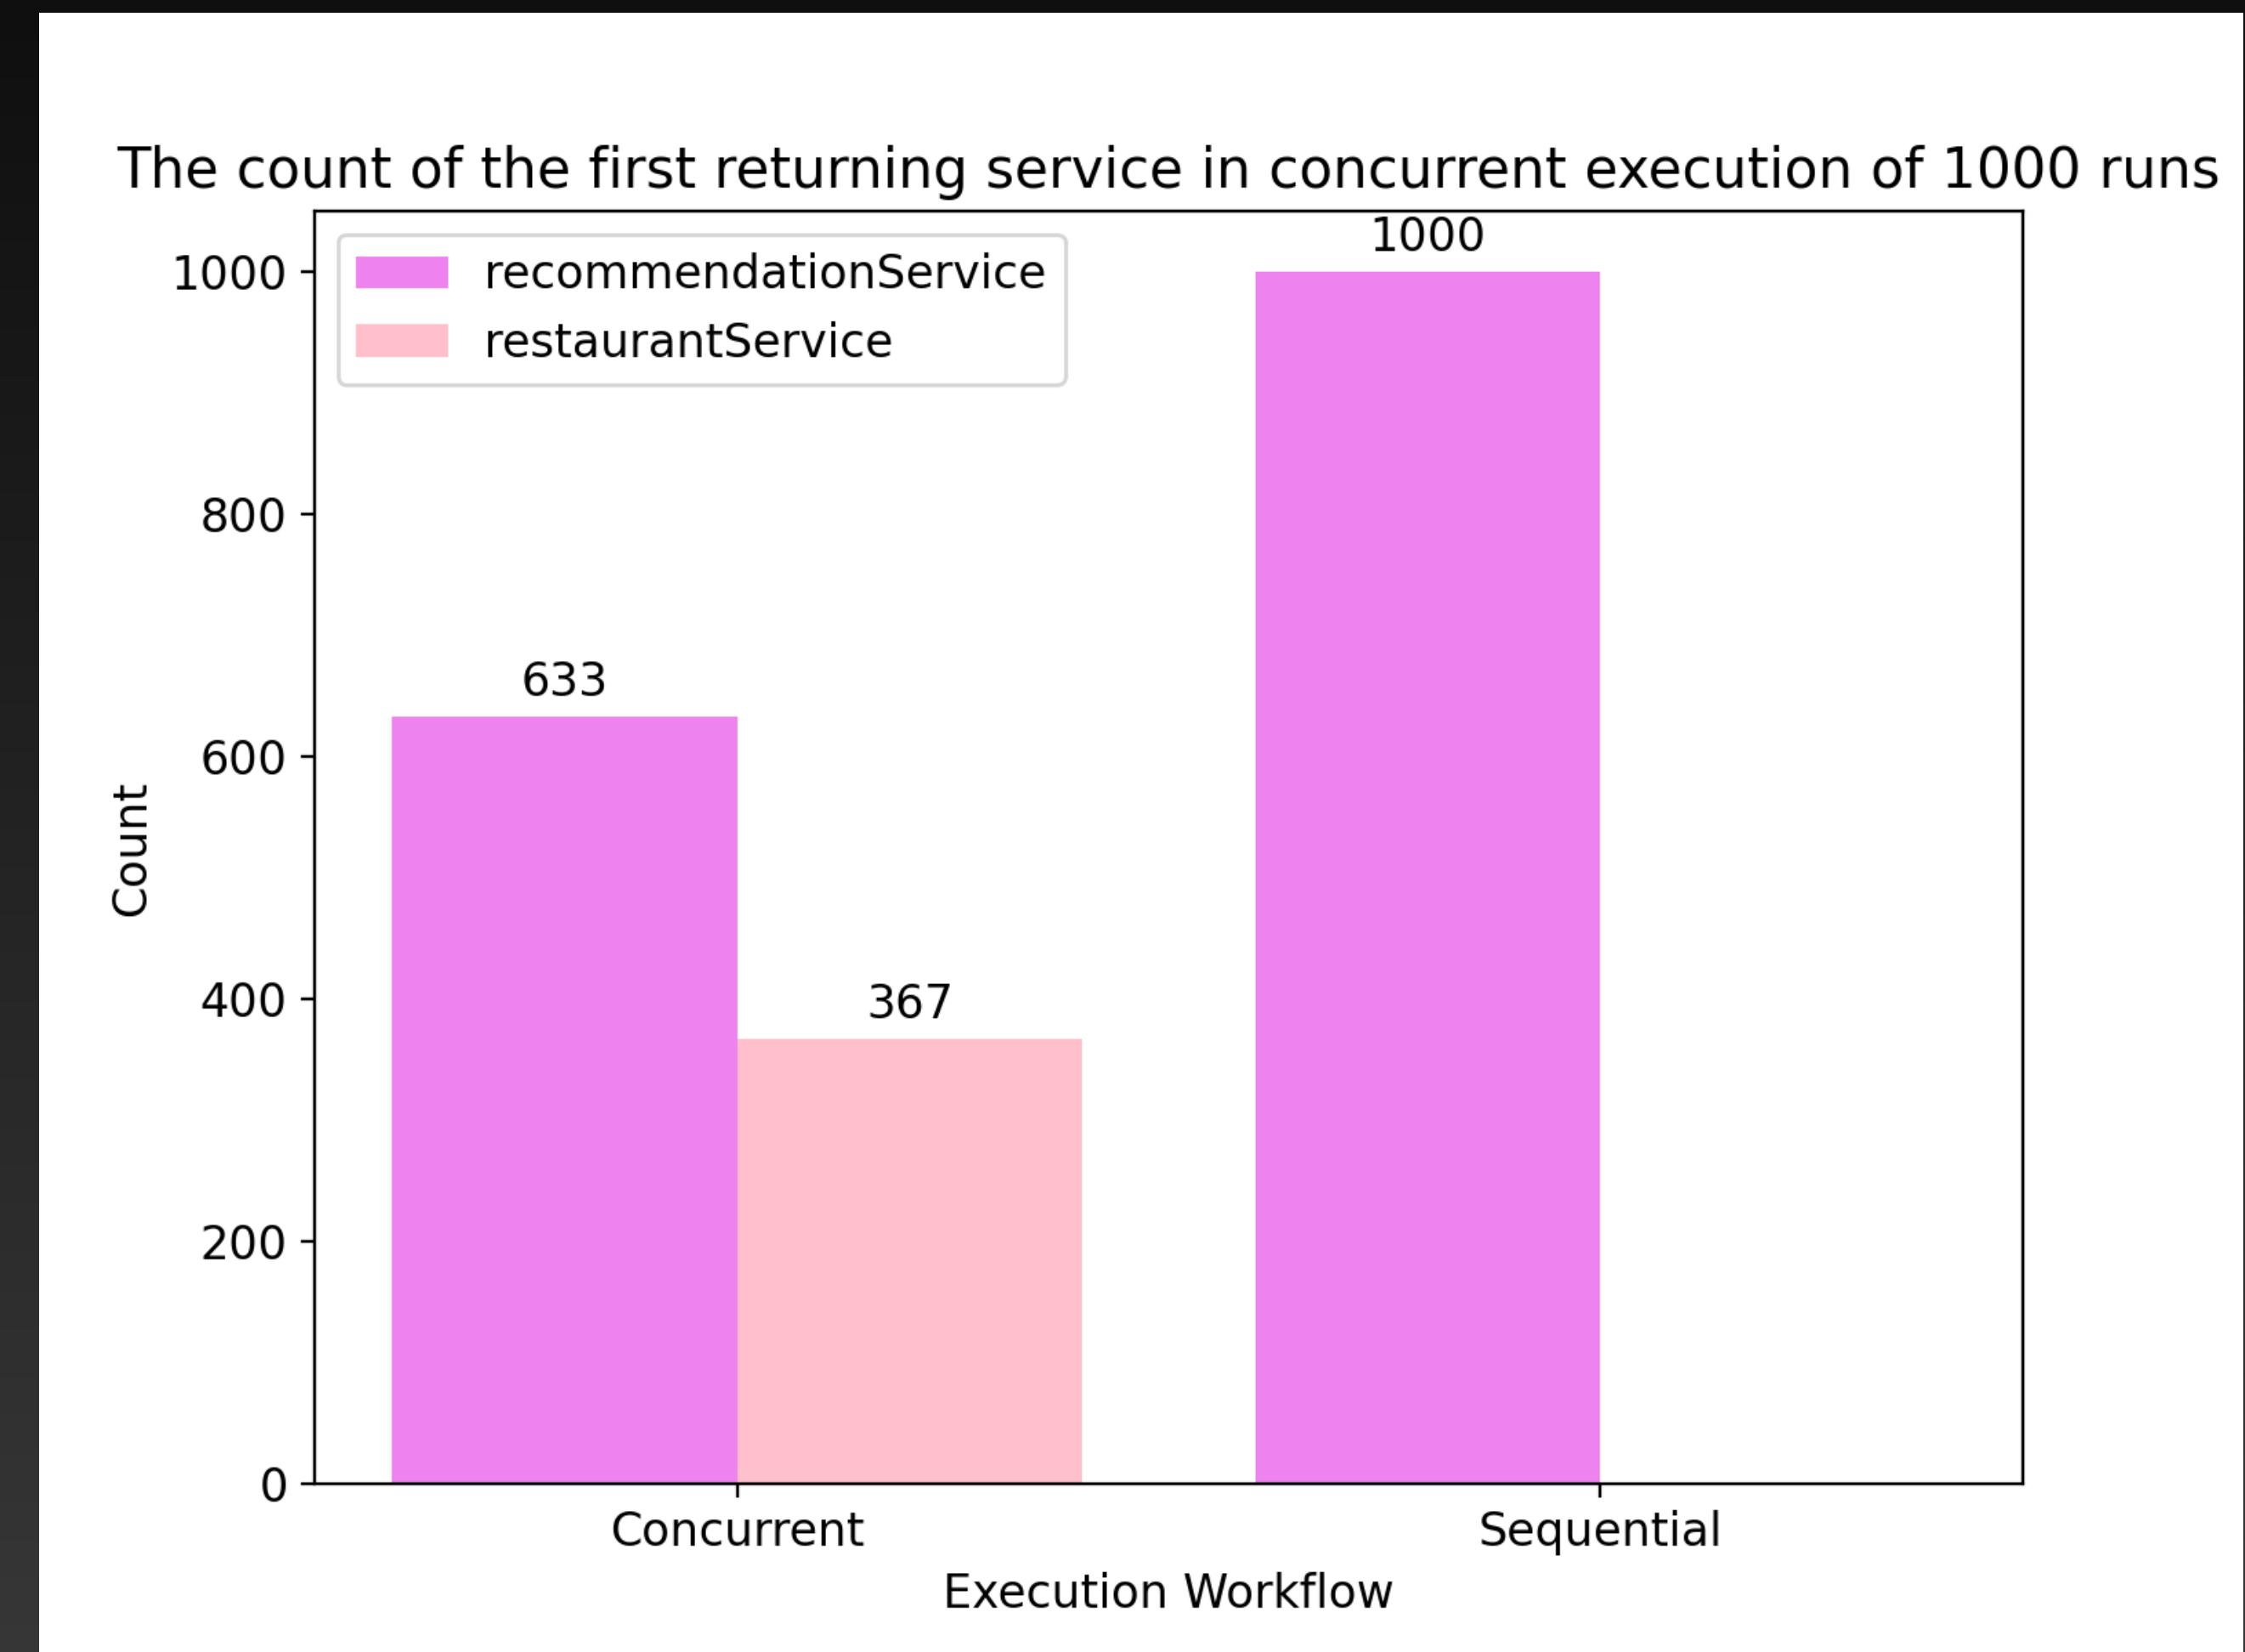


Figure 9: Toy Demo Validation

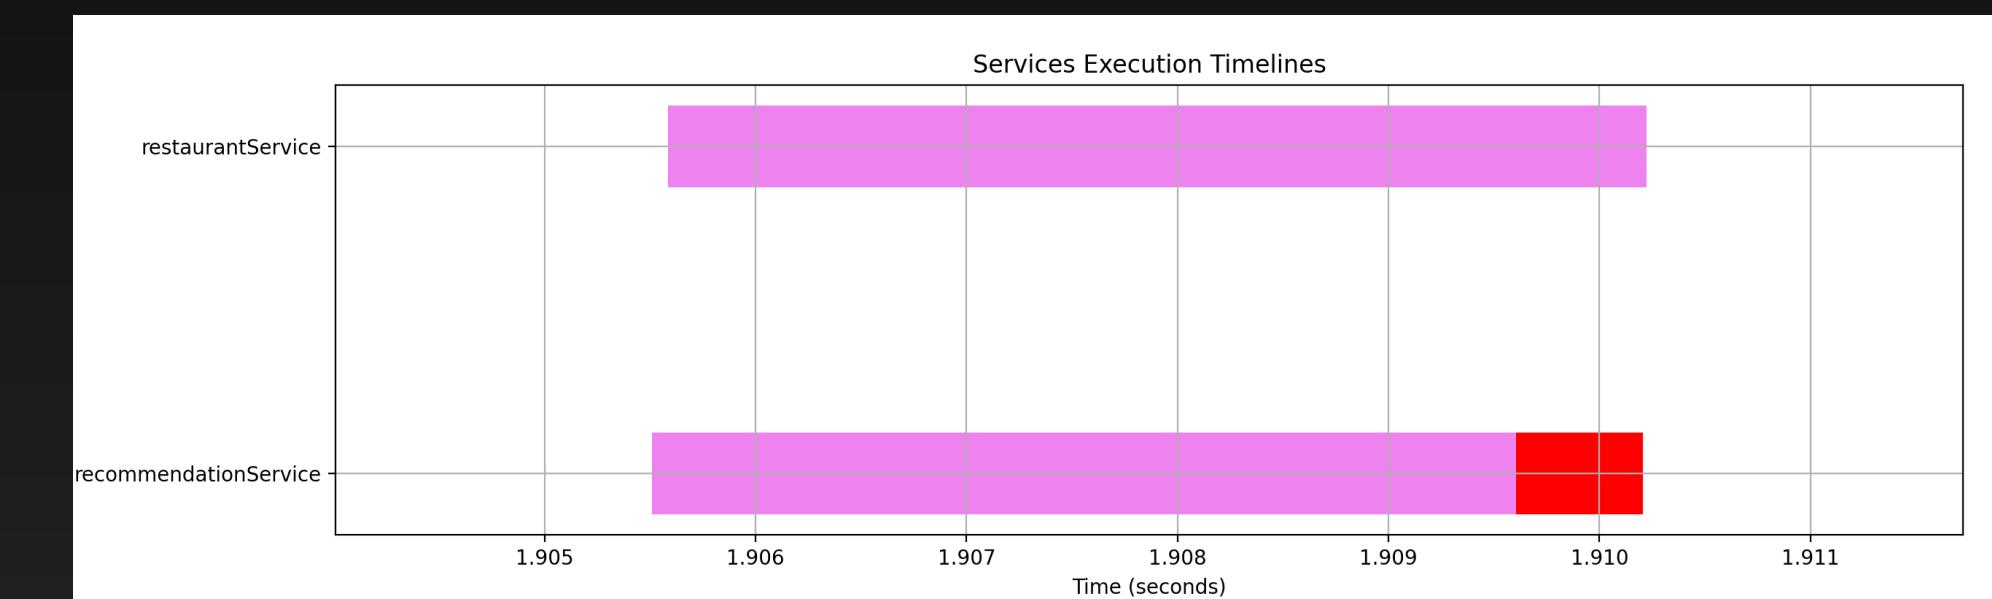
Toy Demo Output

Concurrent

Without Delay



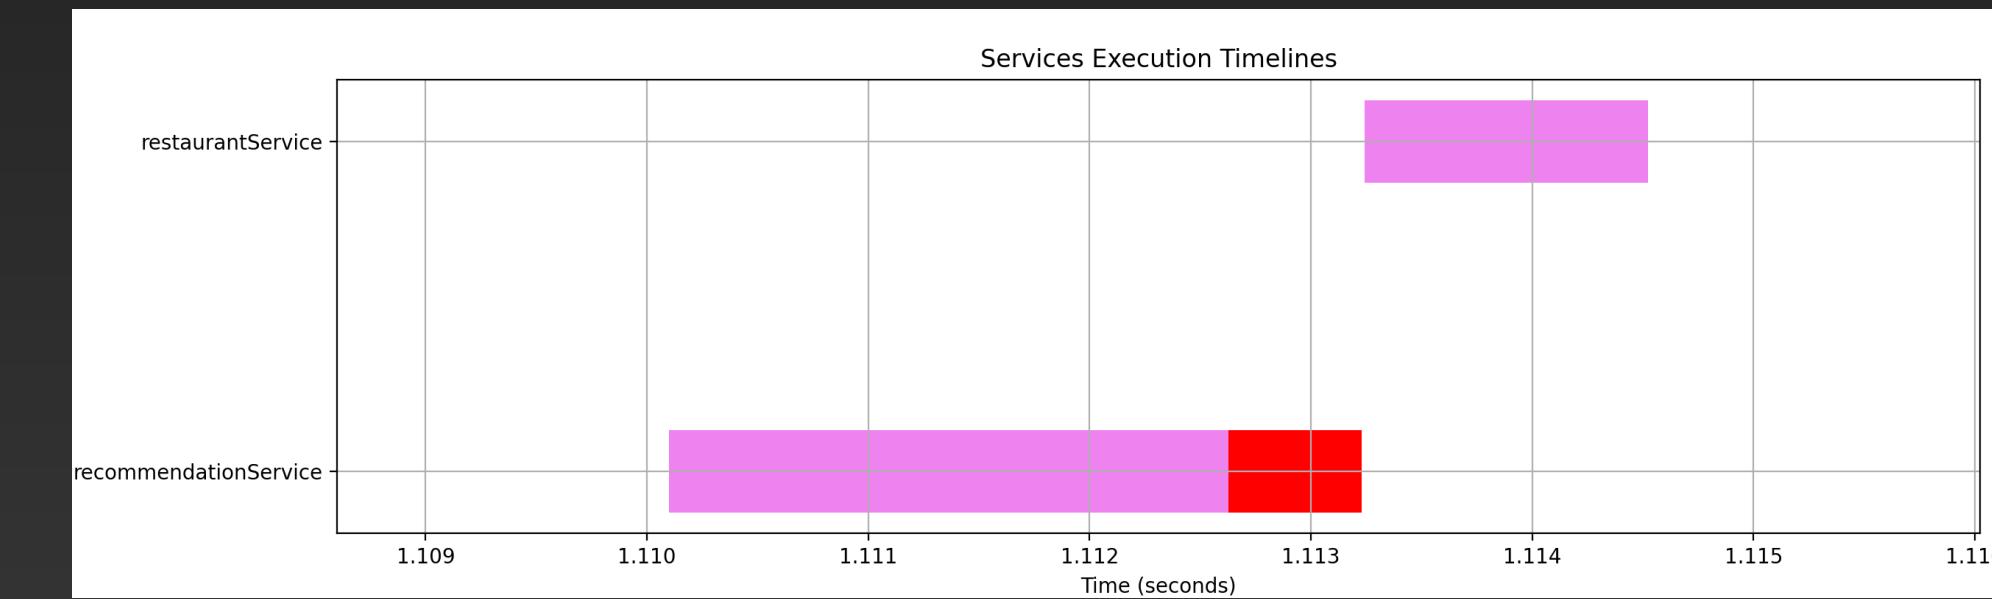
With Delay



Recommendation Service usually returns first (6:4) → only Restaurant Service returns first

Figure 10: Toy Demo Output

Sequential



Recommendation Service always returns first

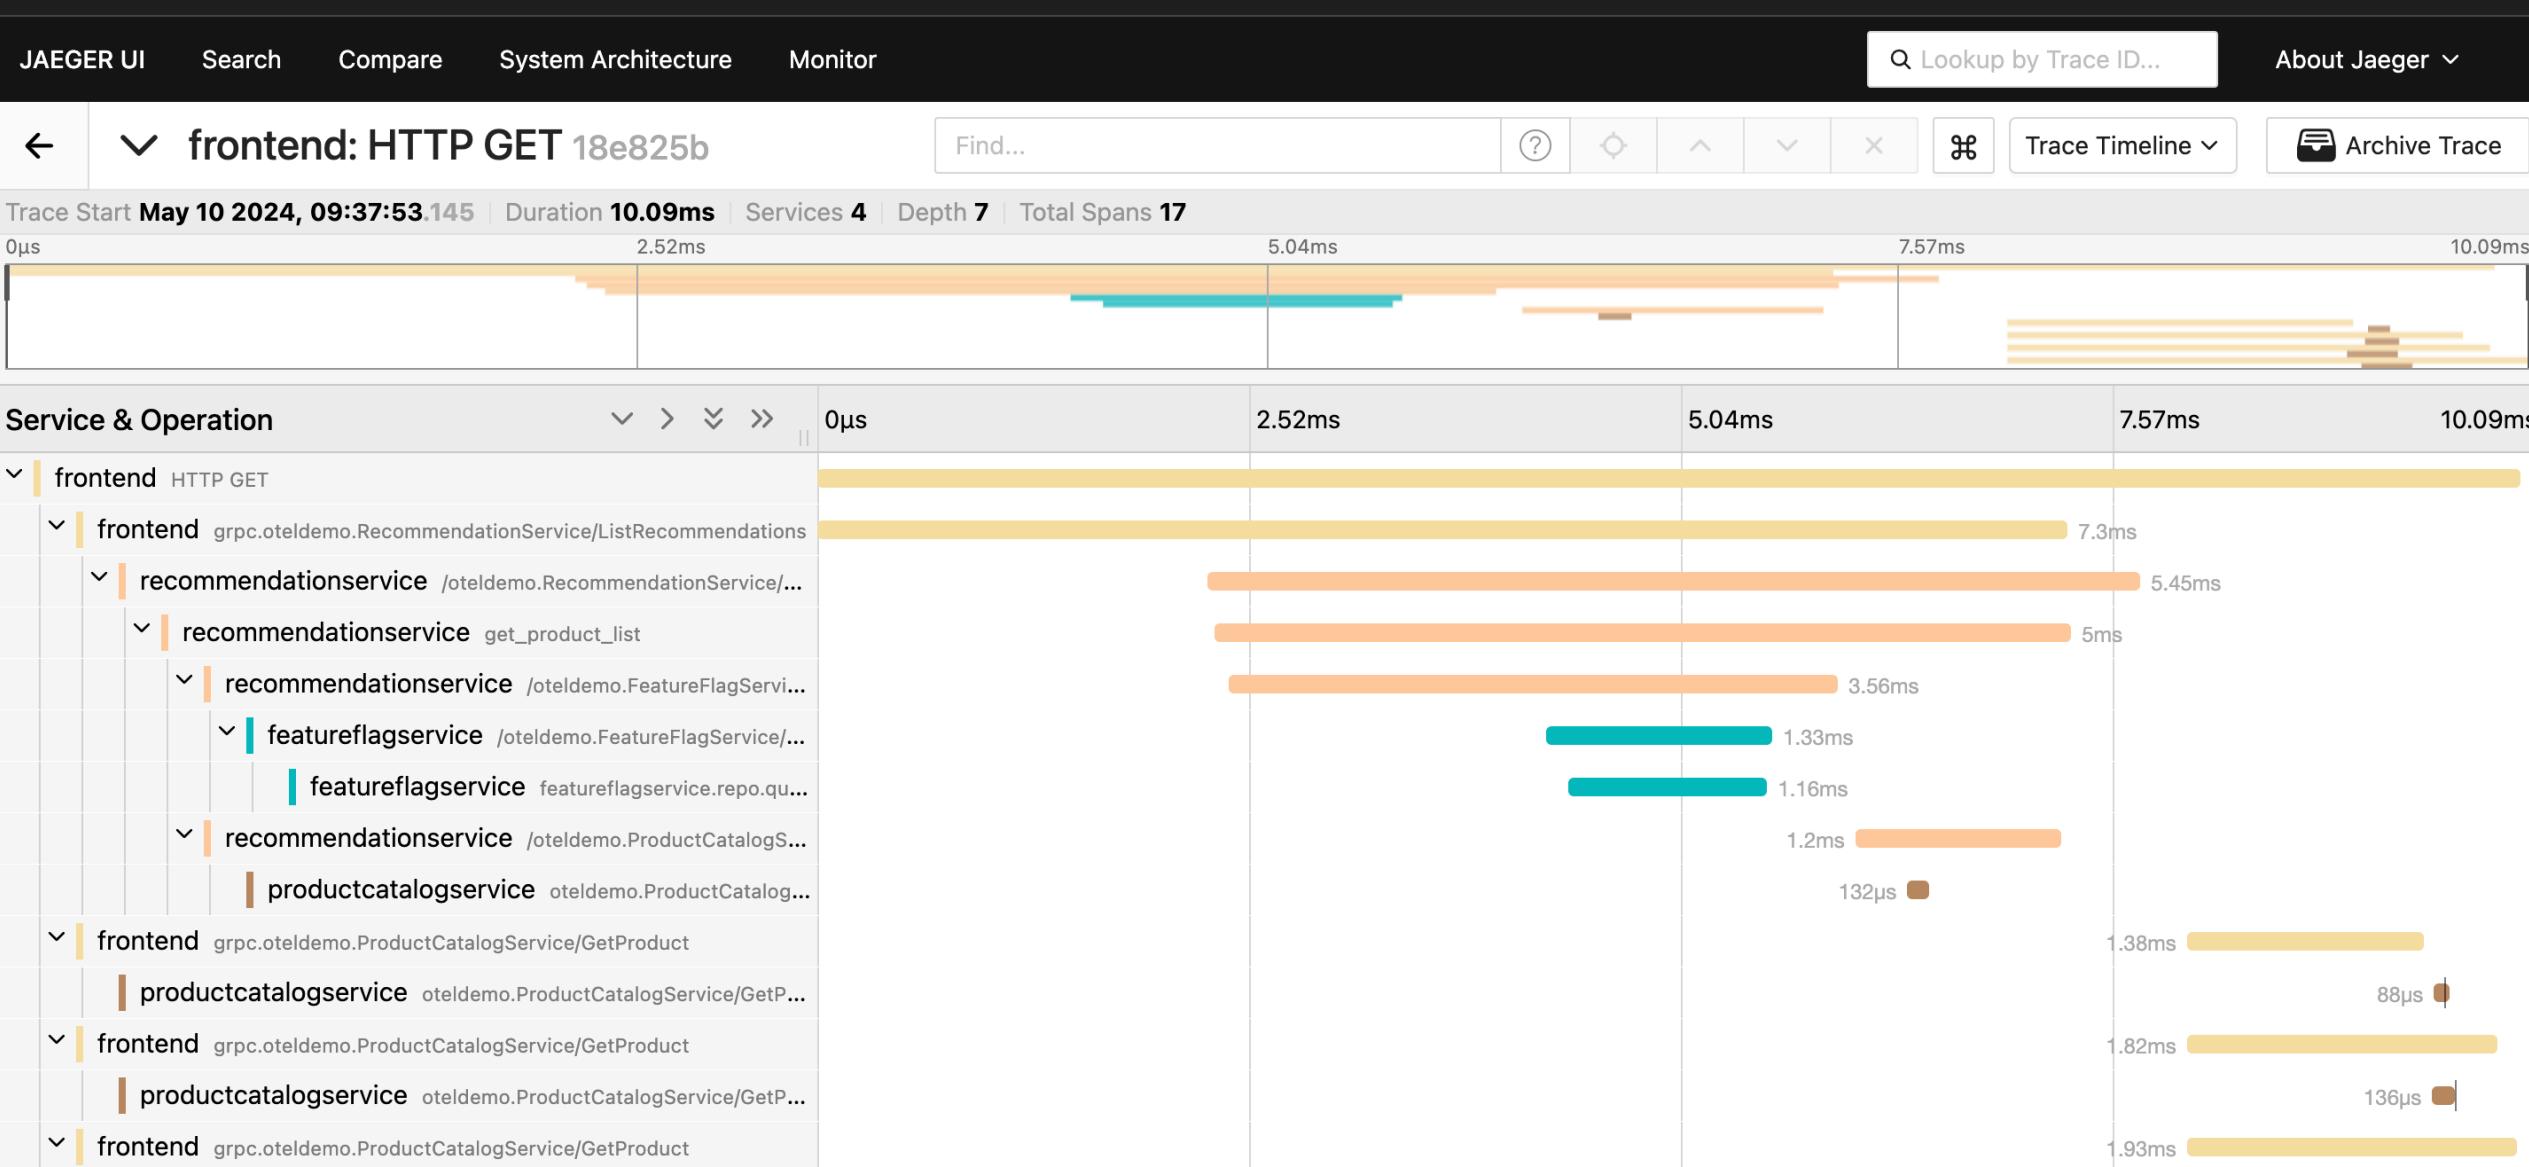
Evaluations

- Add Server Interceptor to OTel Demo
 - Overheads
 - Precision

Is the Result Valuable/Significant?

- Effective dependency analysis
- Construct the true structure -> validated with Jaeger visual
- A working proof of concept

Undisturbed



Disturbed using server-side Interceptor

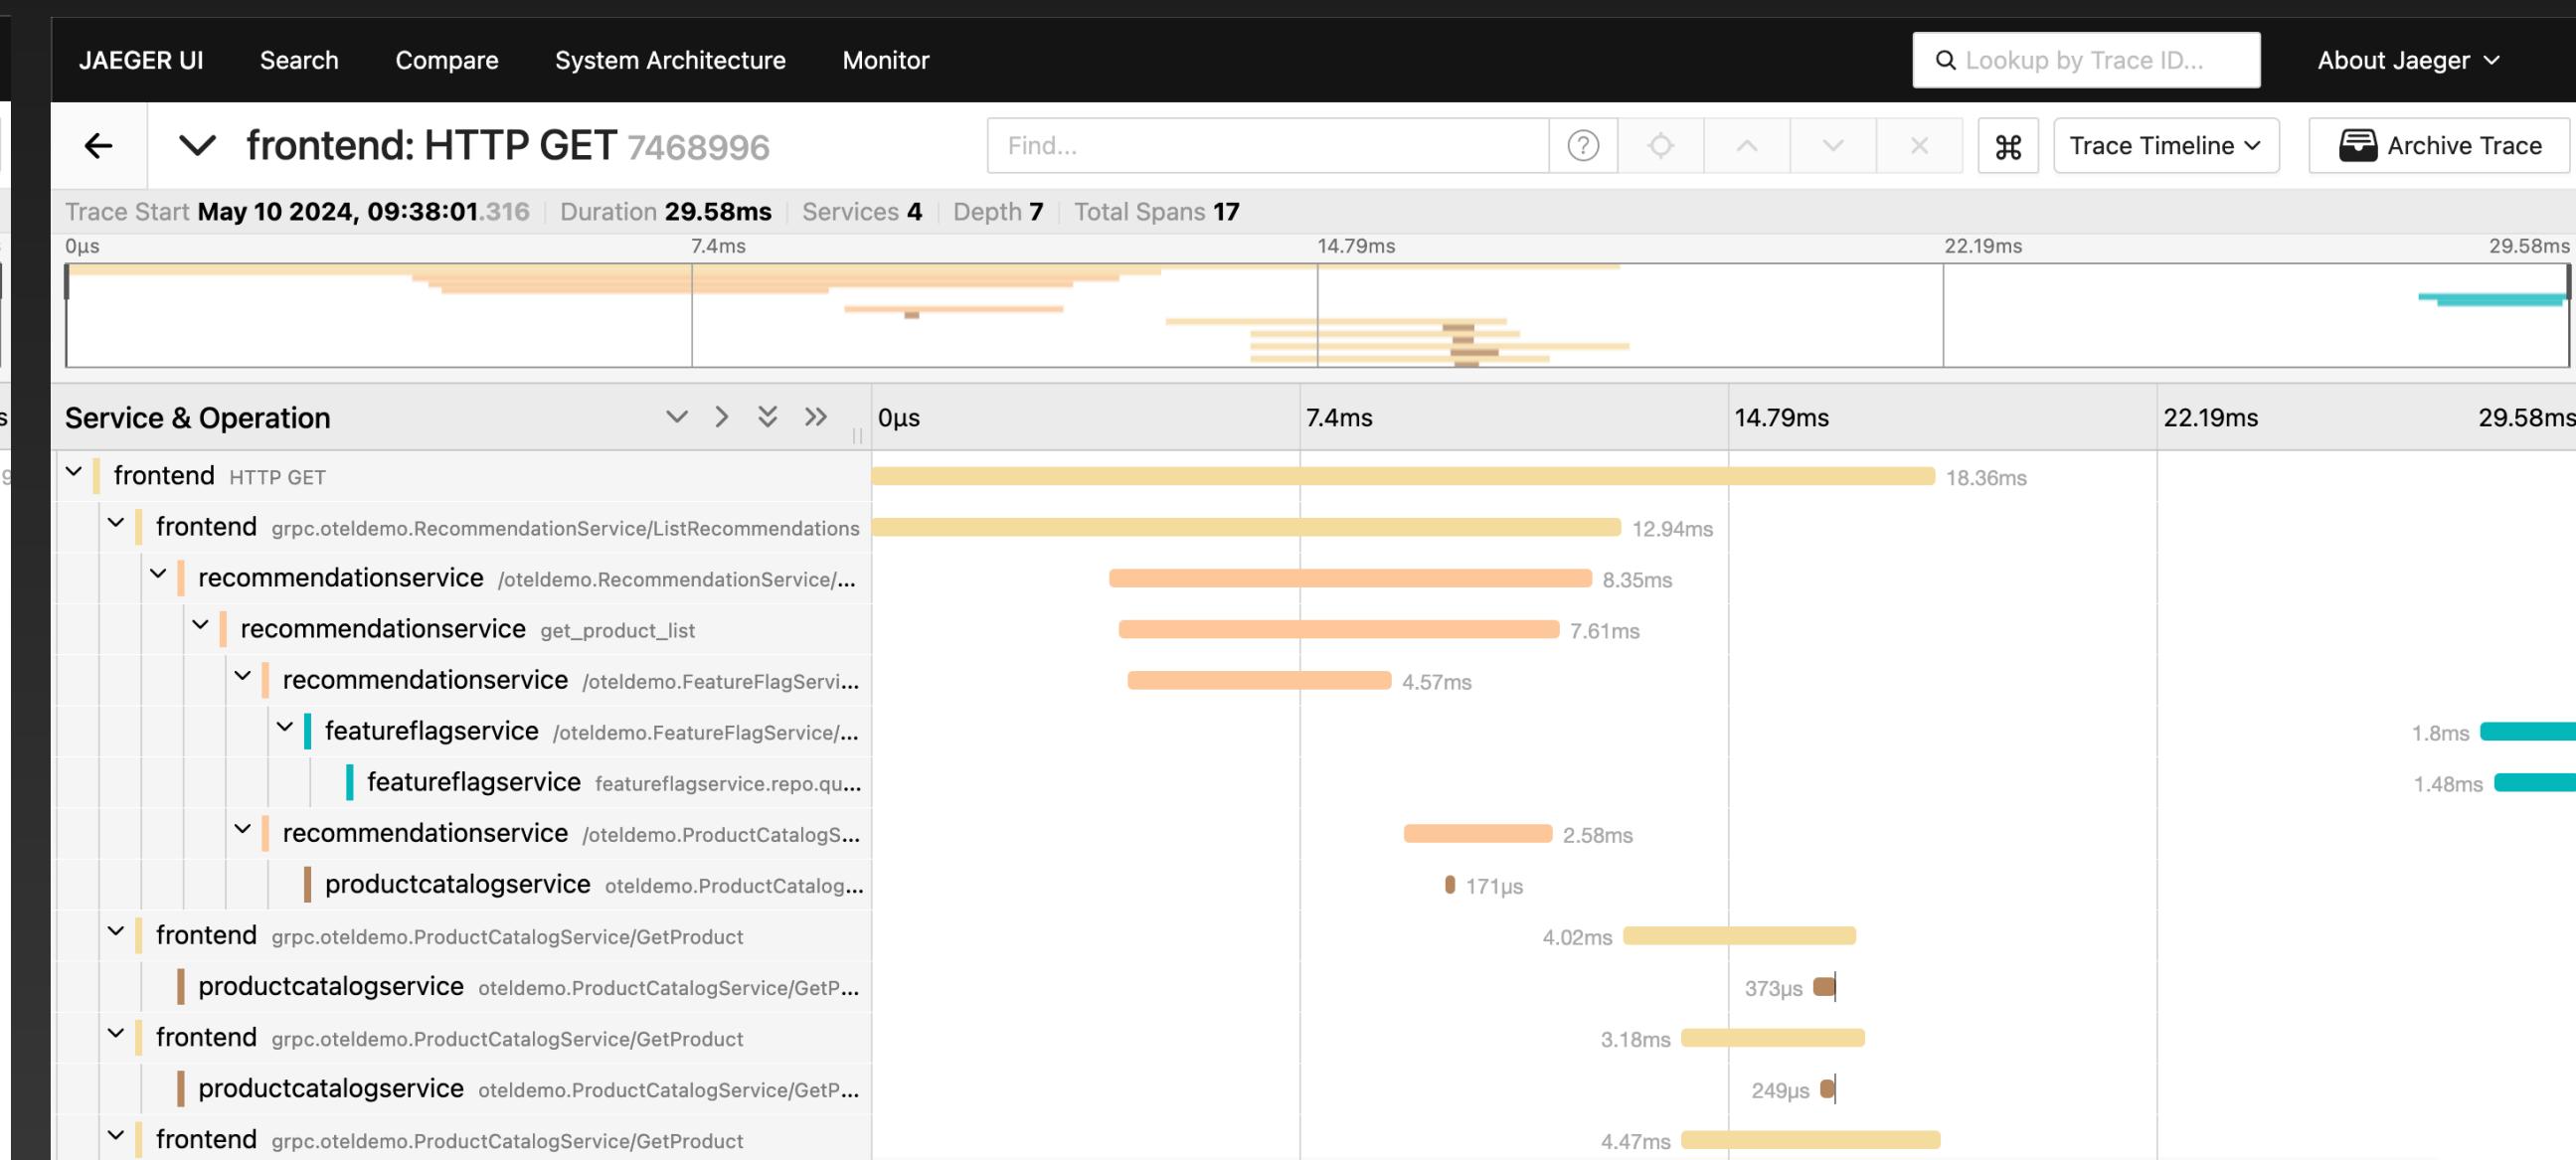


Figure 11: Jeager Trace Comparison

OpenTelemetry Demo

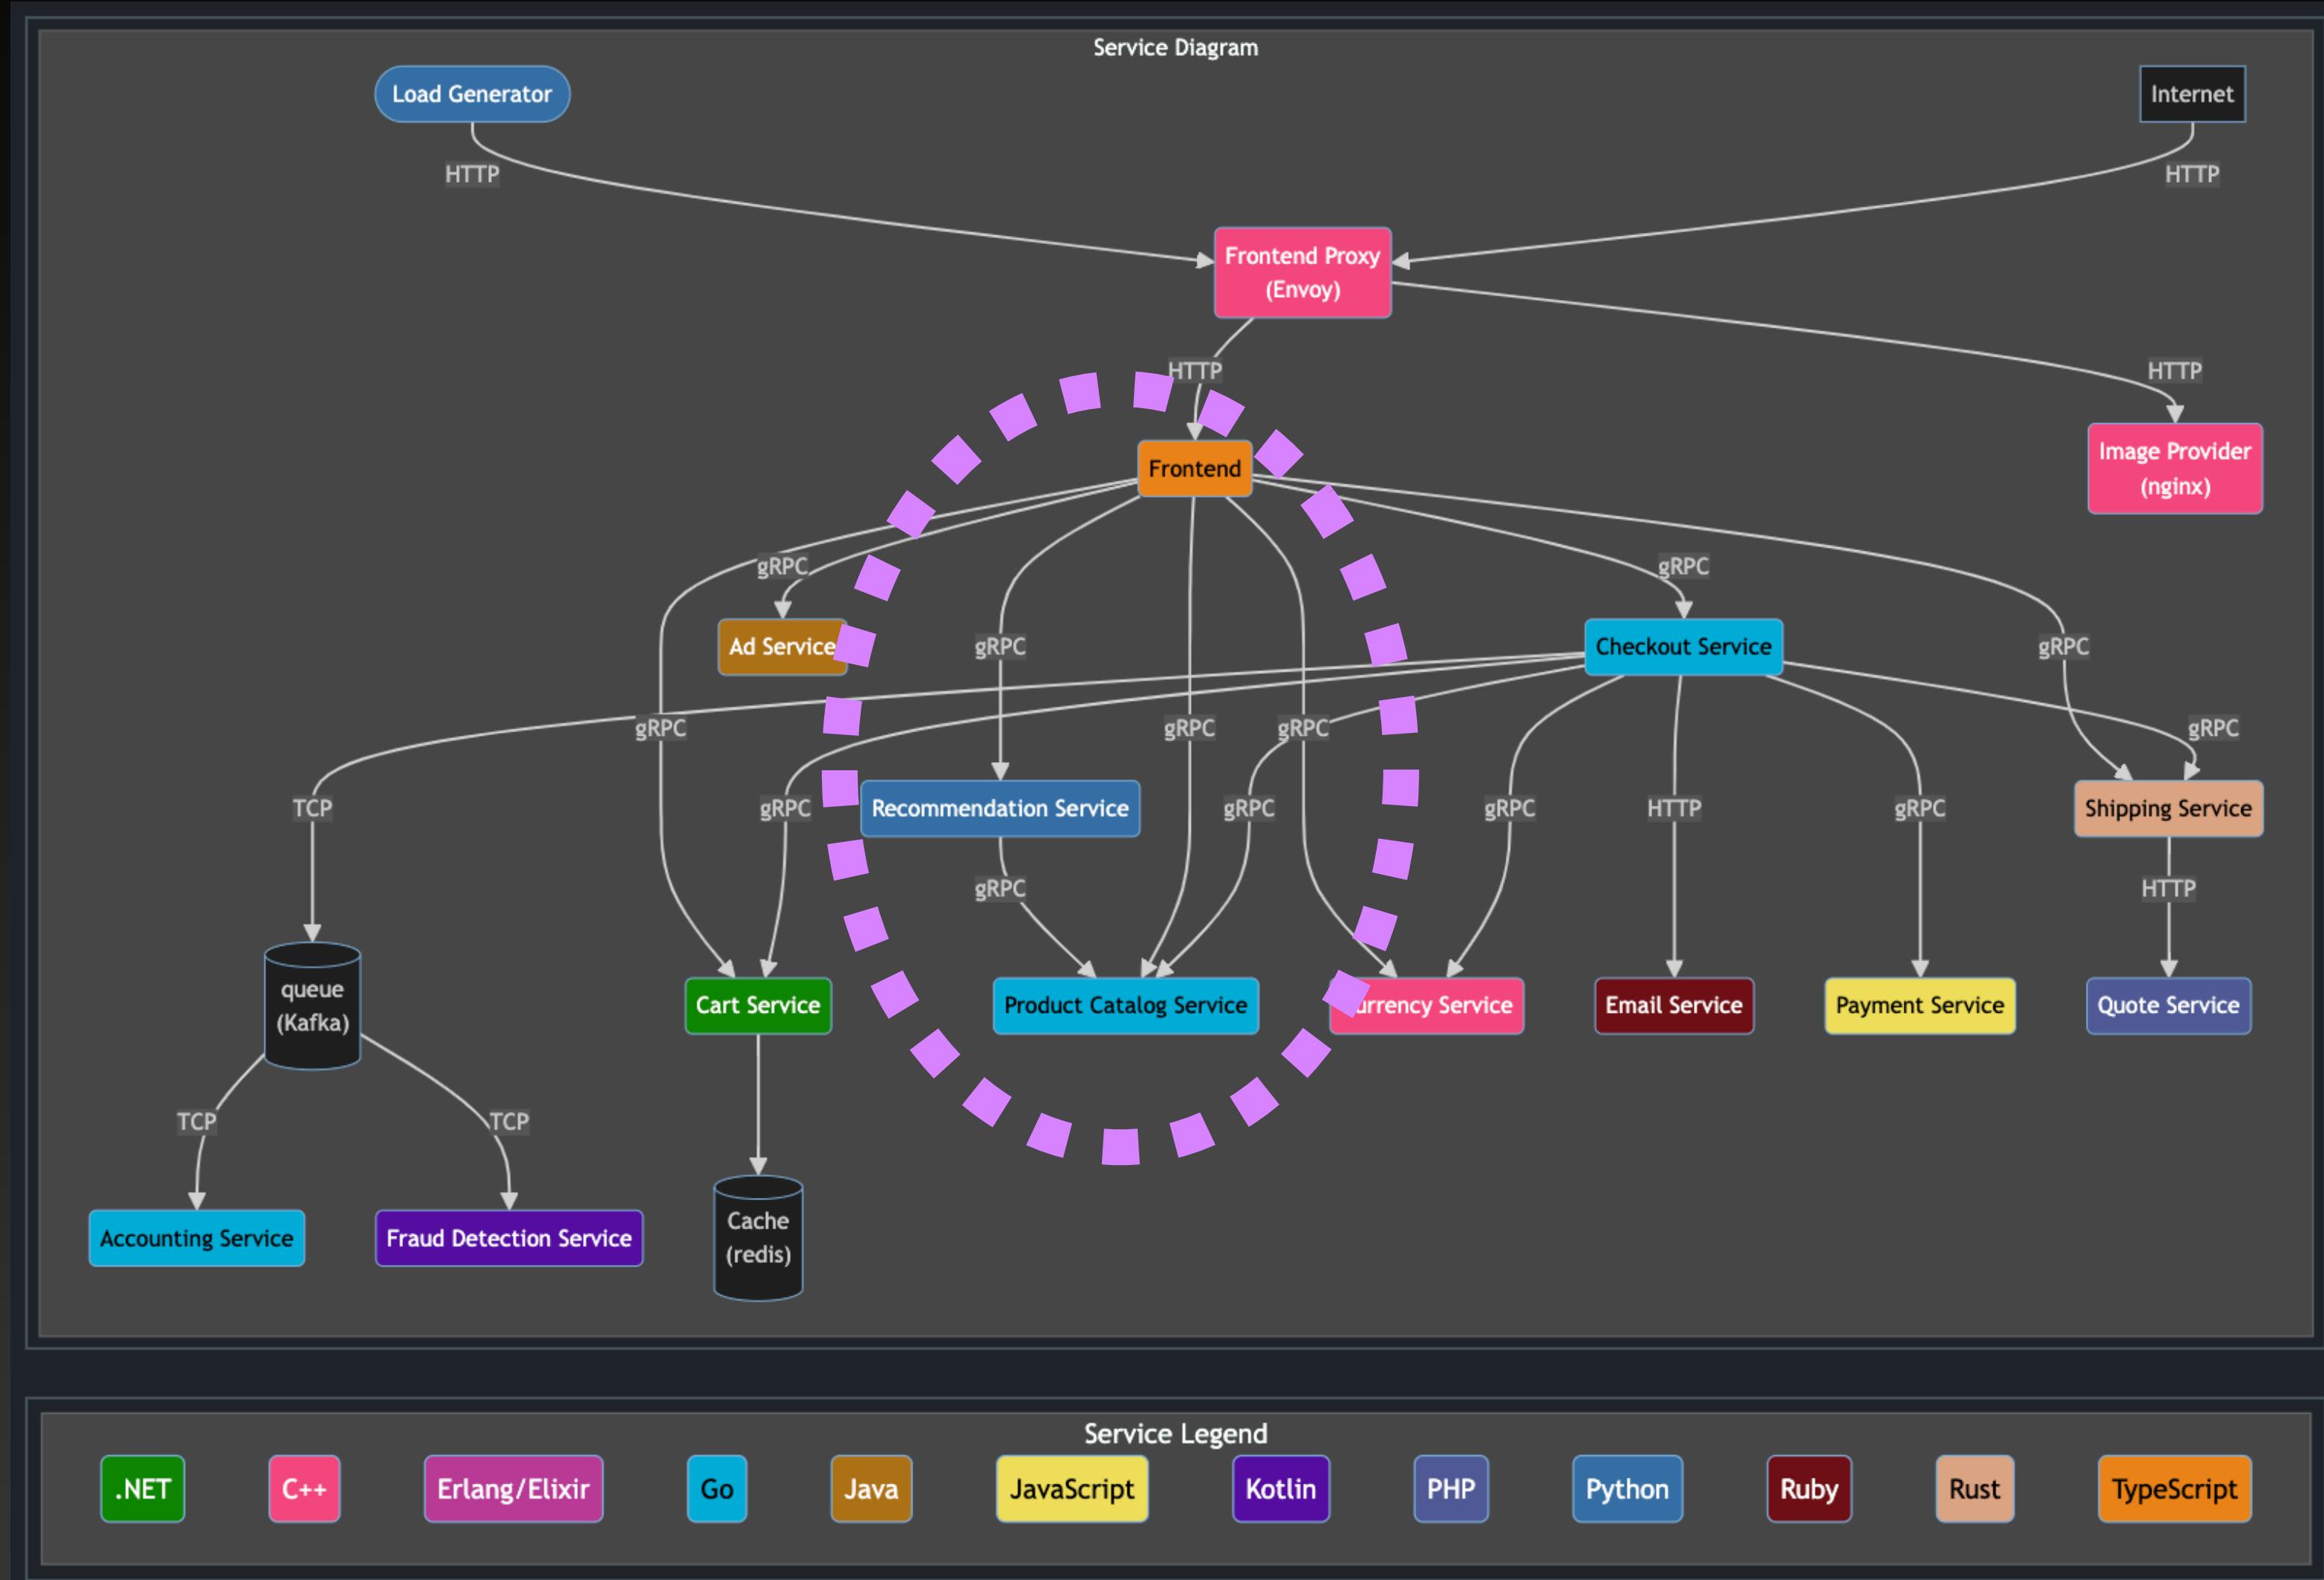


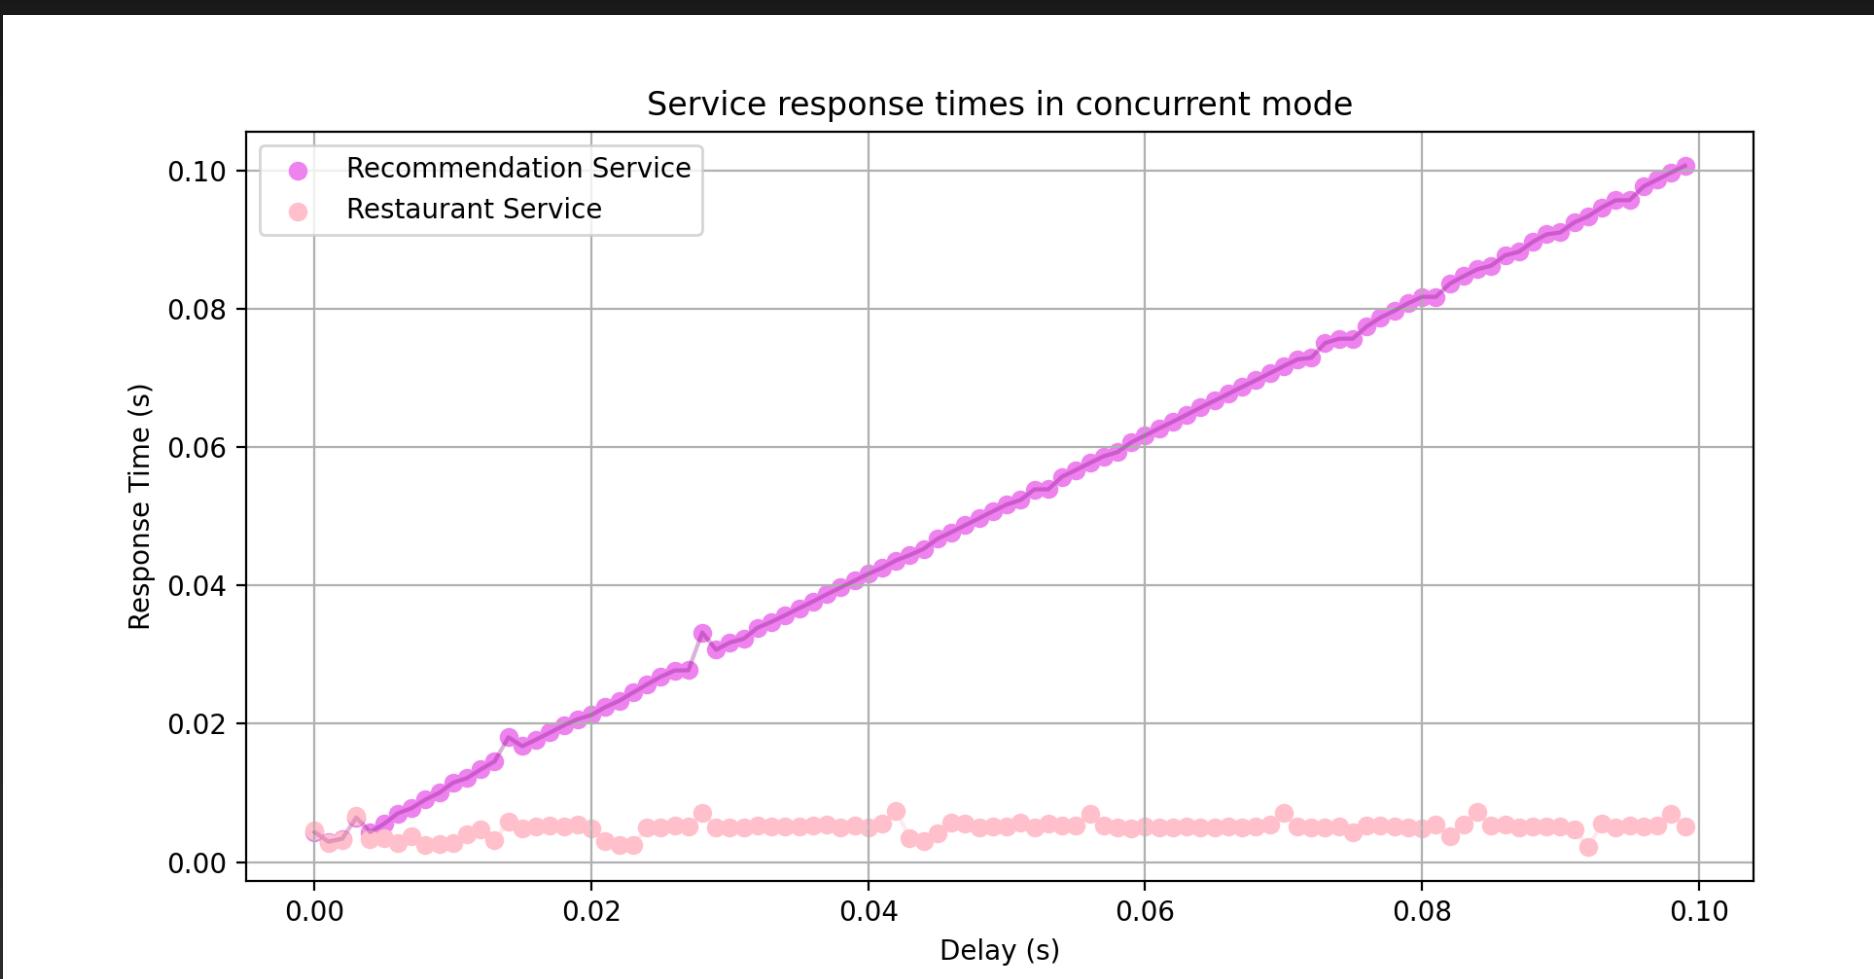
Figure 12: OpenTelemetry Astronomy Shop Microservices Architecture

- 15 lines of code
- Retrieve service name from docker-compose.yml
- 5 seconds delay before returning back to frontend
- Reveal the true dependency: sequential workflow

What is the Sensitivity/Ripple Effect?

- Overheads versus delays (0 to 0.1 with step 0.001)

Concurrent



Sequential

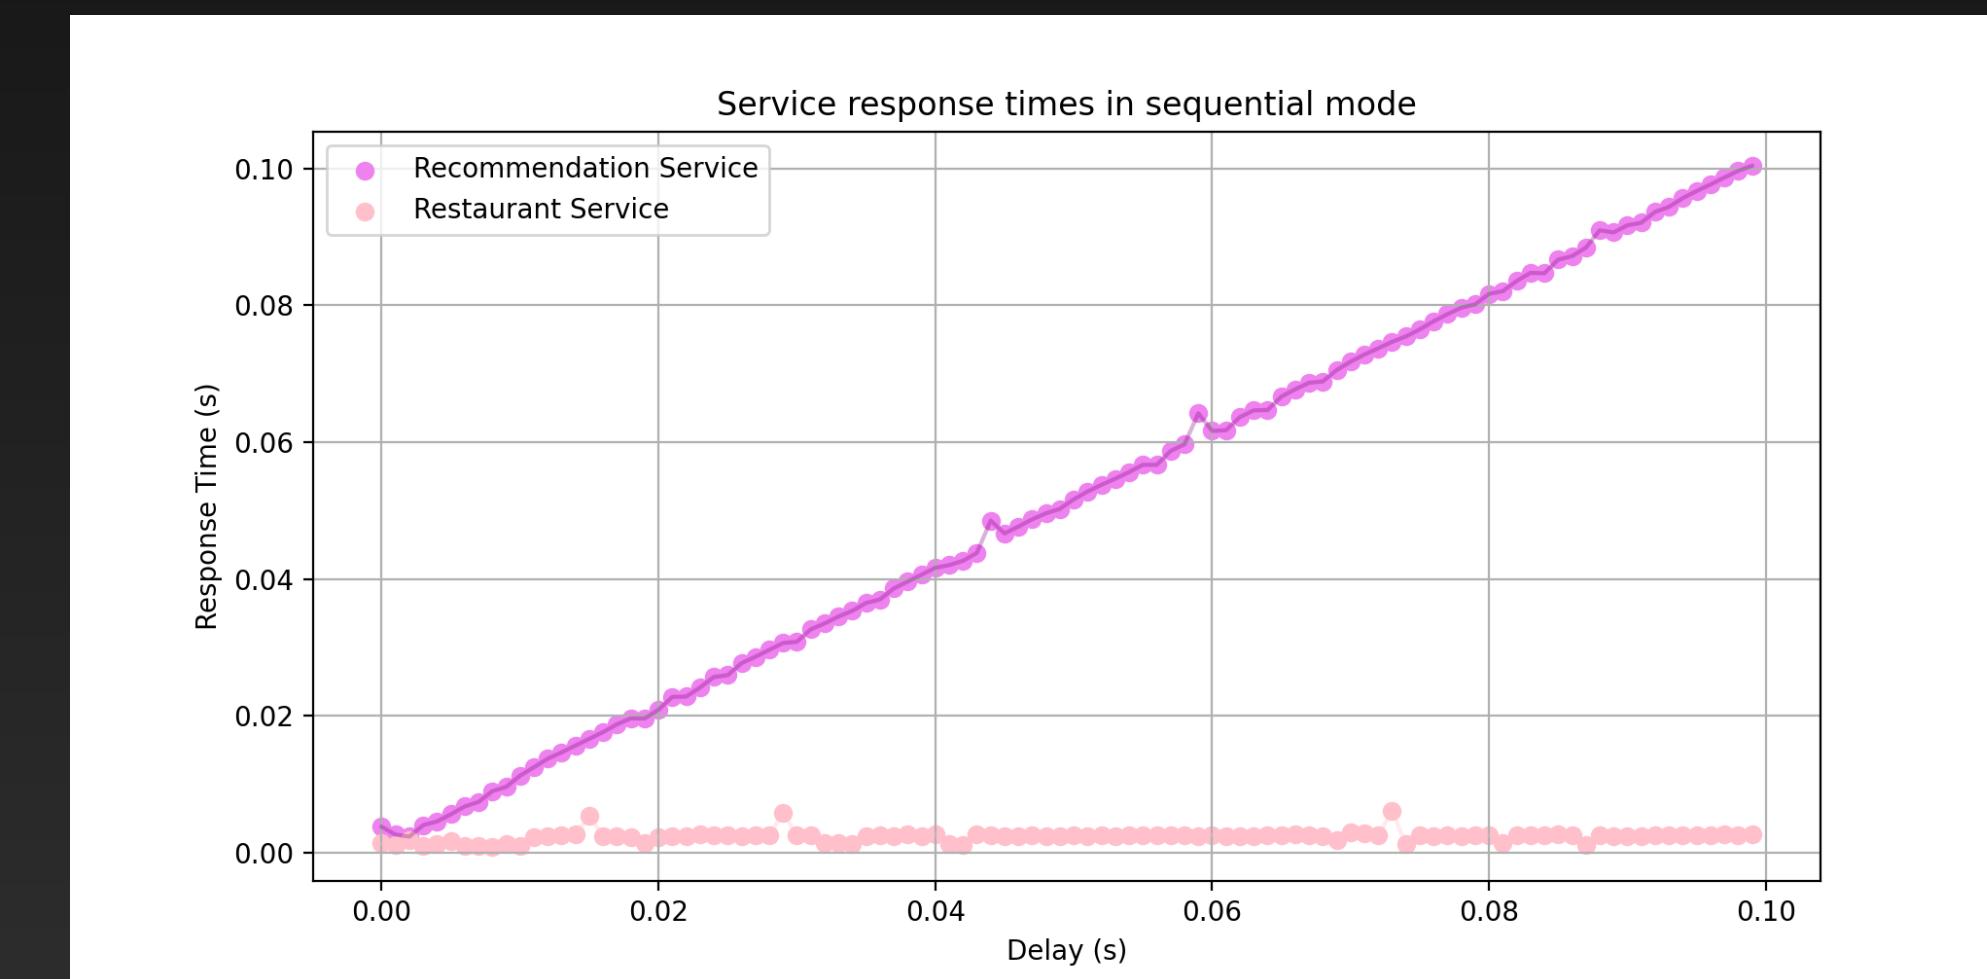


Figure 13: Measuring overheads

- Undisturbed sibling service in concurrent mode experience overhead

What is the Precision: Expected Versus Actual Delay?

- Actual latency is less than intended amount of delay (0.01 seconds)

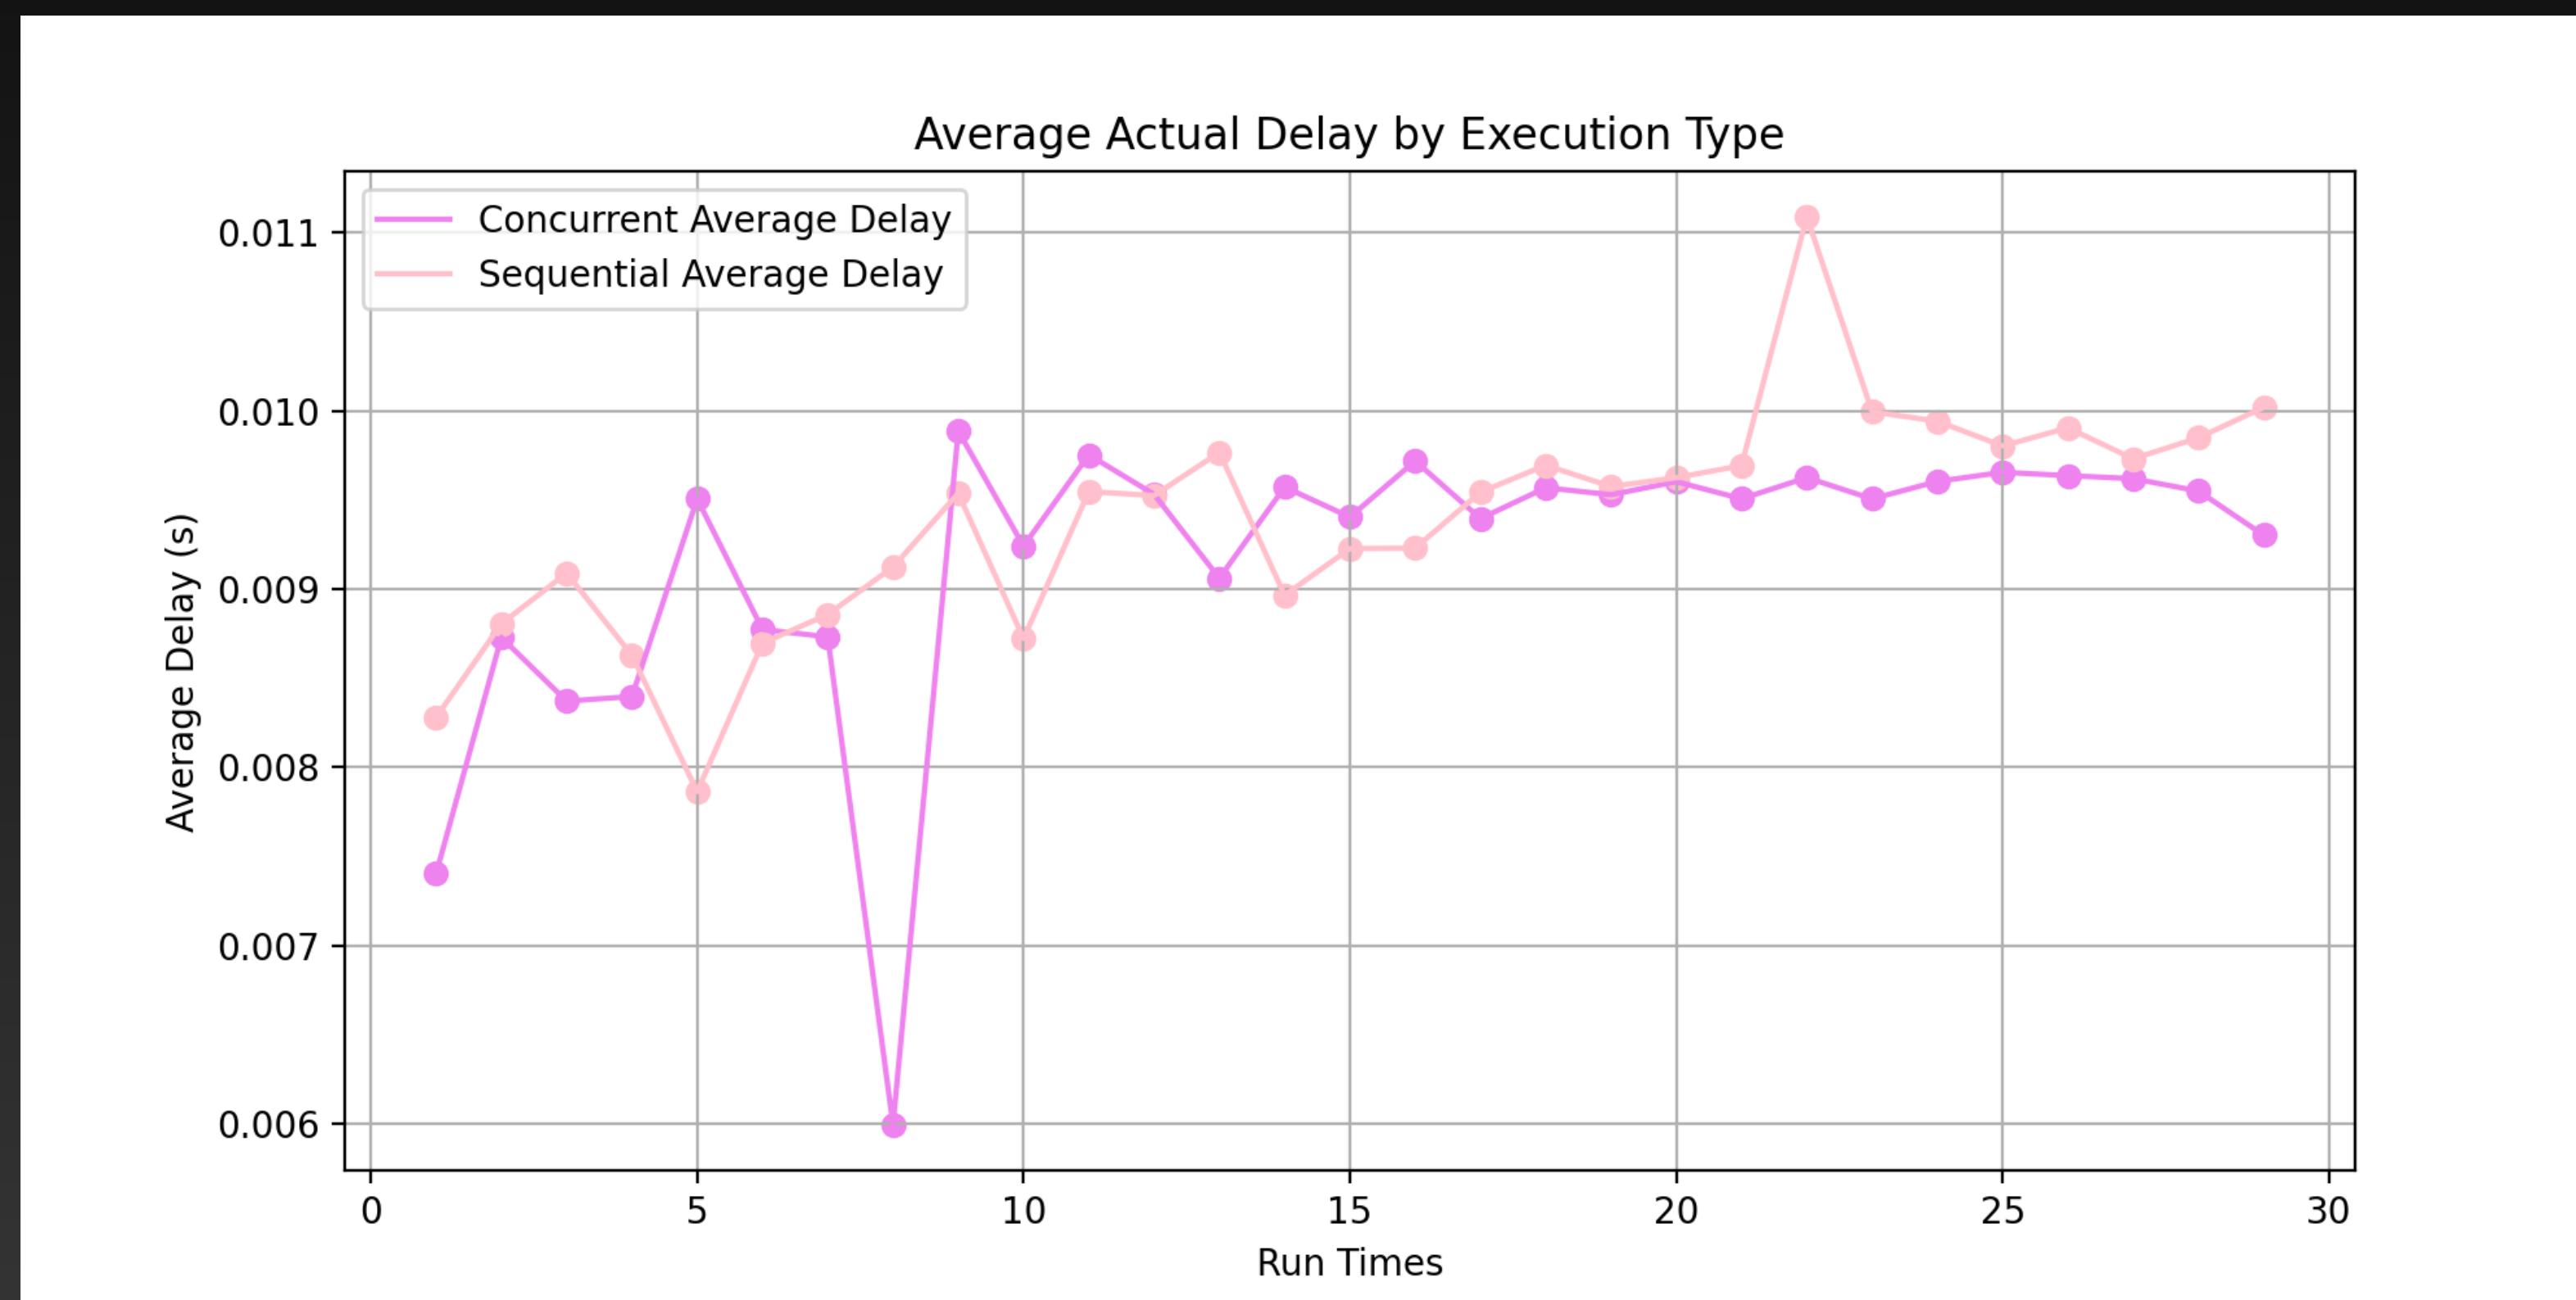


Figure 14: Actual vs expected delays

Future Work

gRPC Interceptors are Limited -> Envoy

Language Support	
Language	Example
C++	C++ example ↗
Go	Go example ↗
Java	Java example ↗
Python	Python example ↗

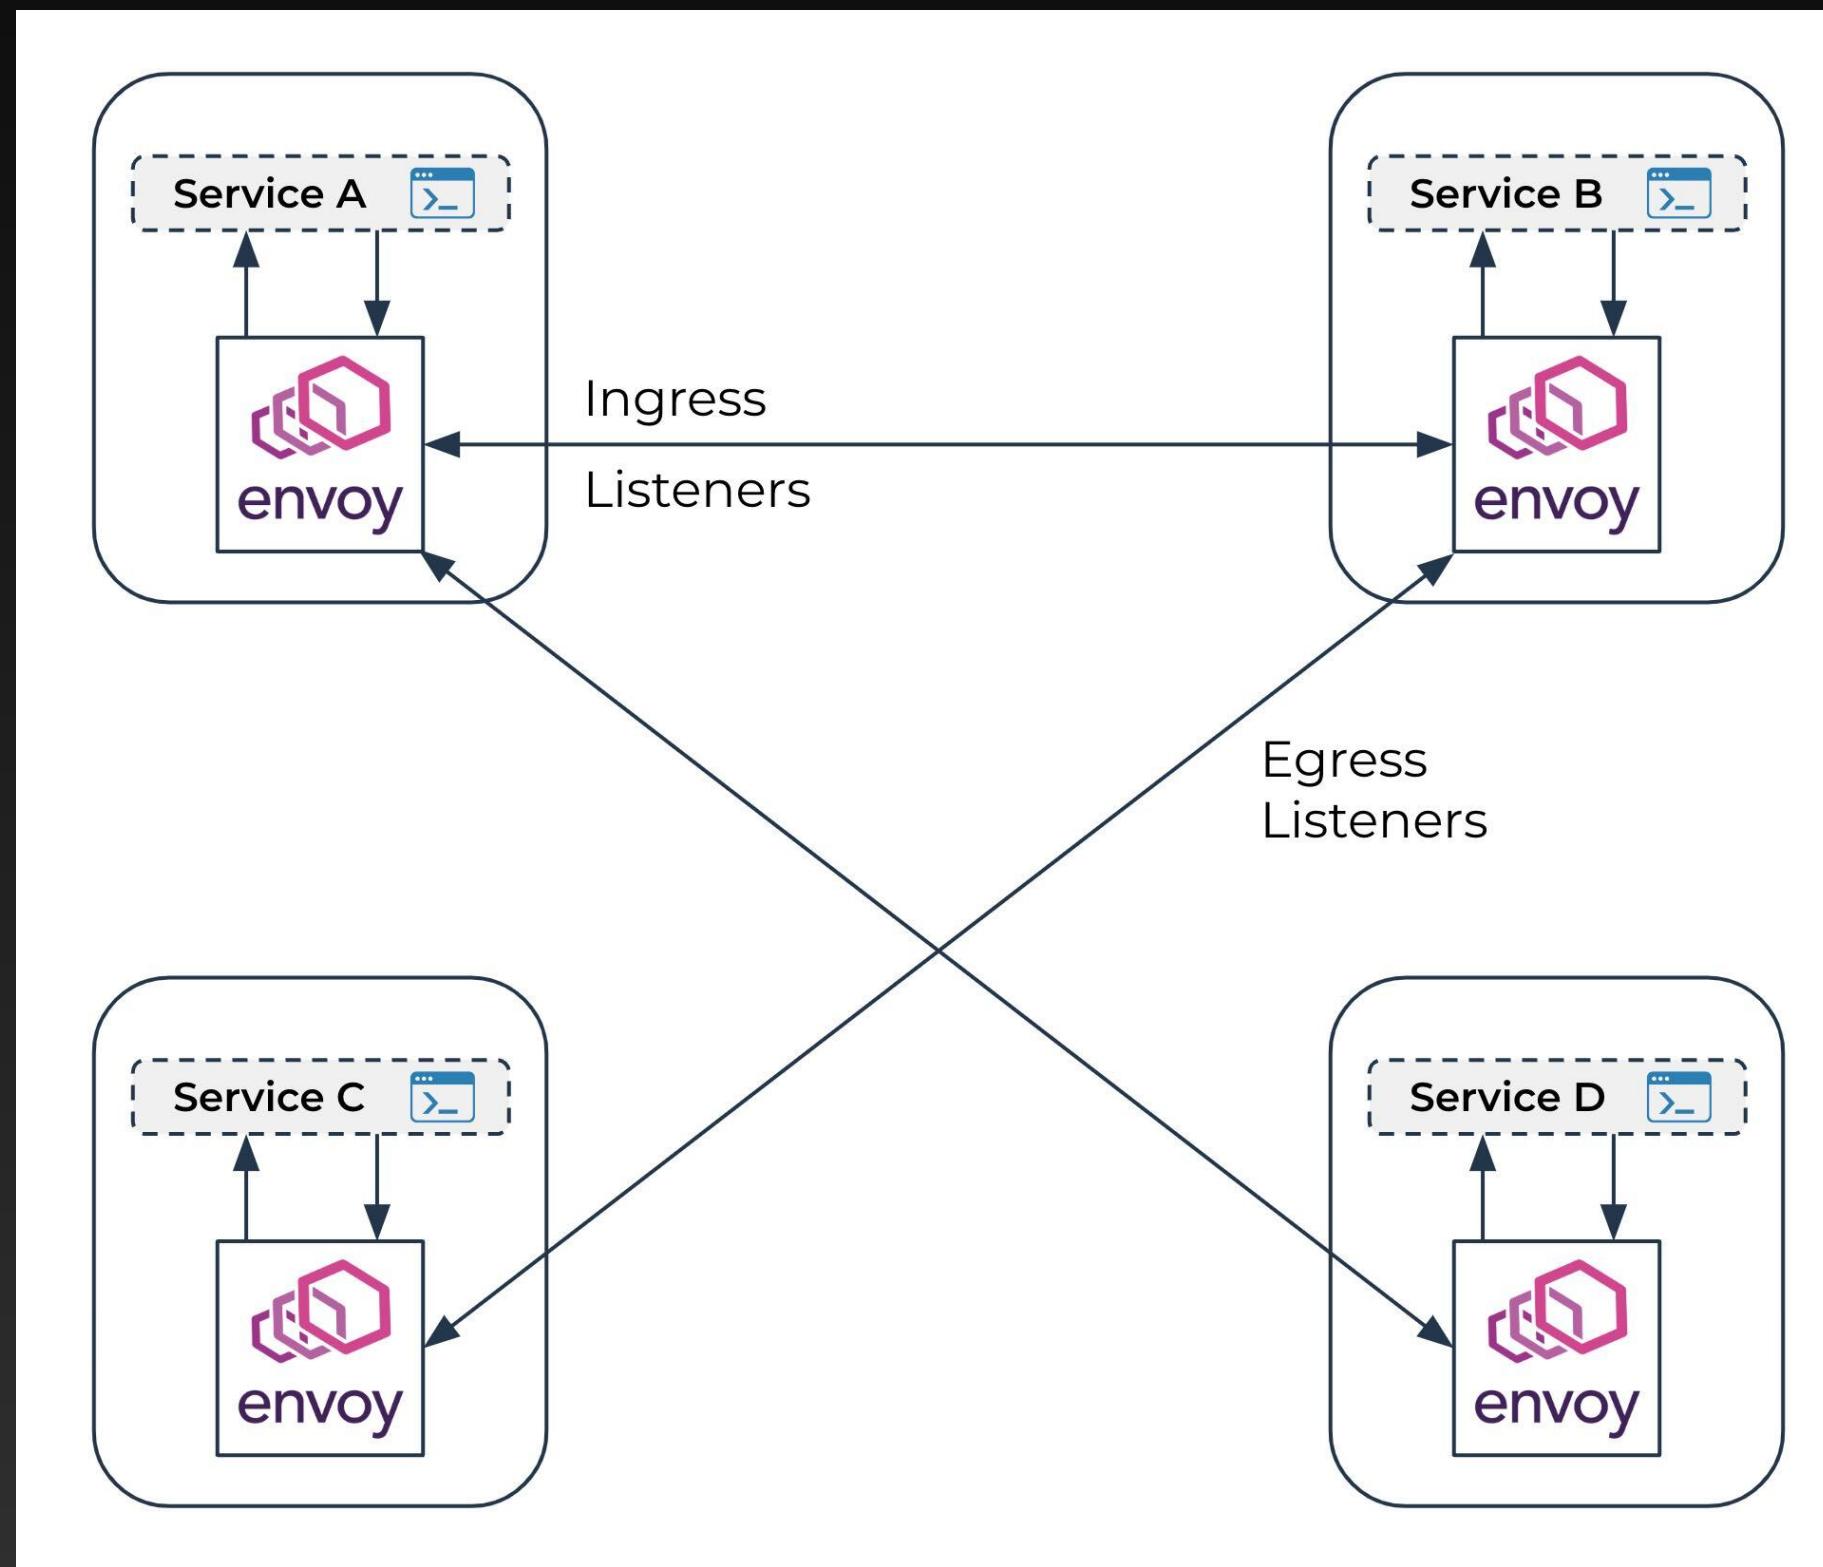


Figure 15: Envoy Architecture

- Open-source service proxy
- Language agnostic
- Observability
- Extensible
- Resource intensive
- Complexity in debugging
- **Hard to compile**

Reference:

Figure 5: <https://www.coingecko.com/learn/crypto-rpc-best-rpc-providers>

Figure 6: <https://grpc.io/docs/what-is-grpc/introduction/>

Figure 12: <https://opentelemetry.io/docs/demo/architecture/>

Figure 15: <https://www.solo.io/topics/envoy-proxy/>

A wide-angle photograph of a dark, star-filled night sky. In the lower half of the image, a range of snow-covered mountains is visible against a dark blue horizon. The upper half of the sky is dominated by the Aurora Borealis, with bright green and yellow-green light streaks and patches illuminating the dark blue void. A single small white star is visible in the upper left quadrant.

Thank you for listening!